

Capstone Project: BAN 693

FORECASTING AMAZON STOCK PRICE

Prepared by:

1. Reena Namani (NetId – AH7791)
2. Bhargav Silaparasetty (NetId – RM2790)

Instructor name: Jiming Wu

Affiliations: California State University - East Bay

Summary:

The objective of the project is to forecast stock price of AMAZON using the historical dataset available, which contains information on Amazon's Stock price, that encompasses the time period from May 1997 to September 2023. The dataset will consist of several columns, each of which offers details regarding its stock performance on a daily basis such as opening price, closing price, day high and day low.

We have used regression analysis, time series analysis, and predictive modeling to analyze the historical dataset of Amazon's stock price and forecast its future values.

We have also used different models such as two-level forecasting, and the advanced Exponential Smoothing & Arima model to improve the accuracy of the forecast. By comparing the RMSE and MAPE values of each model, we have selected the best model for forecasting future stock prices.

The model with the lowest RMSE and MAPE values is generally considered the best fit for forecasting future stock prices.

It is important to note that no forecasting model can provide perfectly accurate predictions, and some degree of error is inevitable. However, by selecting the model with the lowest RMSE and MAPE values, you can minimize the error and increase the likelihood of making informed investment decisions.

The forecasting results can provide valuable insights into the trends and patterns of Amazon's stock performance and help investors and analysts make informed decisions regarding their investment in the stock. Visualizations such as line graphs and charts can also help in understanding the data and communicate the findings effectively.

Introduction

Amazon, one of the biggest and most prosperous technological businesses in the world, is best recognized for their e-commerce & cloud business. Starting out in 1994, it has expanded quickly, entering a variety of new markets and sectors, such as cloud computing, e-commerce & digital content.

The popularity of Amazon's stock, or AMZN, among investors is a result of the firm's solid financial results and its position as a pioneer in the technological sector. The corporation significantly invests in emerging technology, and its advertising division generates most of its revenue. Even though it is prone to market fluctuations, Amazon's stock is regarded as a solid long-term investment because of the company's financial health and dominant market share.

For this project we got the data from Kaggle & we have conducted a thorough analysis of Amazon's stock price using various statistical techniques such as regression analysis, time series analysis, and predictive modeling.

Regression analysis: Regression analysis is a statistical technique used to examine the relationship between two or more variables. In finance and investing, regression analysis is commonly used to analyze the relationship between a stock's price and other factors that may influence it, such as company earnings, industry trends, or macroeconomic indicators. Regression analysis can help investors and analysts understand how changes in these factors are likely to impact the stock's price.

Time series analysis: Time series analysis is a statistical technique used to analyze data that is collected over time, such as stock prices or economic indicators. The goal of time series analysis is to identify patterns or trends in the data and make predictions about future values based on those patterns. Time series analysis involves techniques such as trend analysis, seasonal analysis, and forecasting.

Predictive modeling: Predictive modeling is a statistical technique used to make predictions about future events based on historical data. In finance and investing, predictive modeling is commonly used to forecast future stock prices or other financial metrics such as earnings or revenue. Predictive modeling involves using statistical algorithms and techniques to analyze historical data and identify patterns or trends that can be used to make predictions about the future.

We have also used different models such as two-level forecasting, advanced Exponential Smoothing, and ARIMA models can be used to improve the accuracy of forecasting future stock prices.

Two-Level Forecasting: Two-level forecasting is a statistical technique that uses a combination of regression analysis and time series analysis to forecast future values. The technique involves first using regression analysis to model the relationship between the stock price and external factors that may influence it, such as earnings, industry trends, or macroeconomic indicators. Next, the time series analysis is used to model the time-dependent factors that may impact the stock price. By combining these two models, two-level forecasting can provide a more accurate forecast of future stock prices.

Advanced Exponential Smoothing: Exponential smoothing is a time series analysis technique used to model time-dependent data, such as stock prices. Advanced Exponential Smoothing techniques build on this basic method by introducing additional parameters that can help capture more complex patterns and trends in the data. For example, Holt's Exponential Smoothing adds a trend component to the model, while Winter's Exponential Smoothing adds a seasonal component. By adjusting these parameters, advanced Exponential Smoothing techniques can provide more accurate forecasts of future stock prices.

ARIMA Models: ARIMA (Autoregressive Integrated Moving Average) models are another type of time series analysis technique that can be used to model the time-dependent factors that impact stock prices. ARIMA models are particularly useful when the data exhibits trends or seasonality. ARIMA models involve selecting the appropriate parameters for the autoregressive

(AR), integrated (I), and moving average (MA) components of the model. By selecting the appropriate parameters and fitting the model to historical data, ARIMA models can provide more accurate forecasts of future stock prices.

After creating these models, the accuracy of each model is assessed by comparing the Root Mean Squared Error (RMSE) and Mean Absolute Percentage Error (MAPE) values of each model. RMSE measures the difference between predicted and actual values, while MAPE measures the percentage difference between predicted and actual values. The model with the lowest RMSE and MAPE values is generally considered the best fit for forecasting future stock prices.

In summary, different models such as two-level forecasting, advanced Exponential Smoothing, and ARIMA models can be used to improve the accuracy of forecasting future stock prices. By comparing the RMSE and MAPE values of each model, we have selected the best model for forecasting future stock prices.

Eight steps of forecasting

The above summary describes a project for forecasting the stock price of Amazon using historical data. Let's apply the eight steps of forecasting to this project:

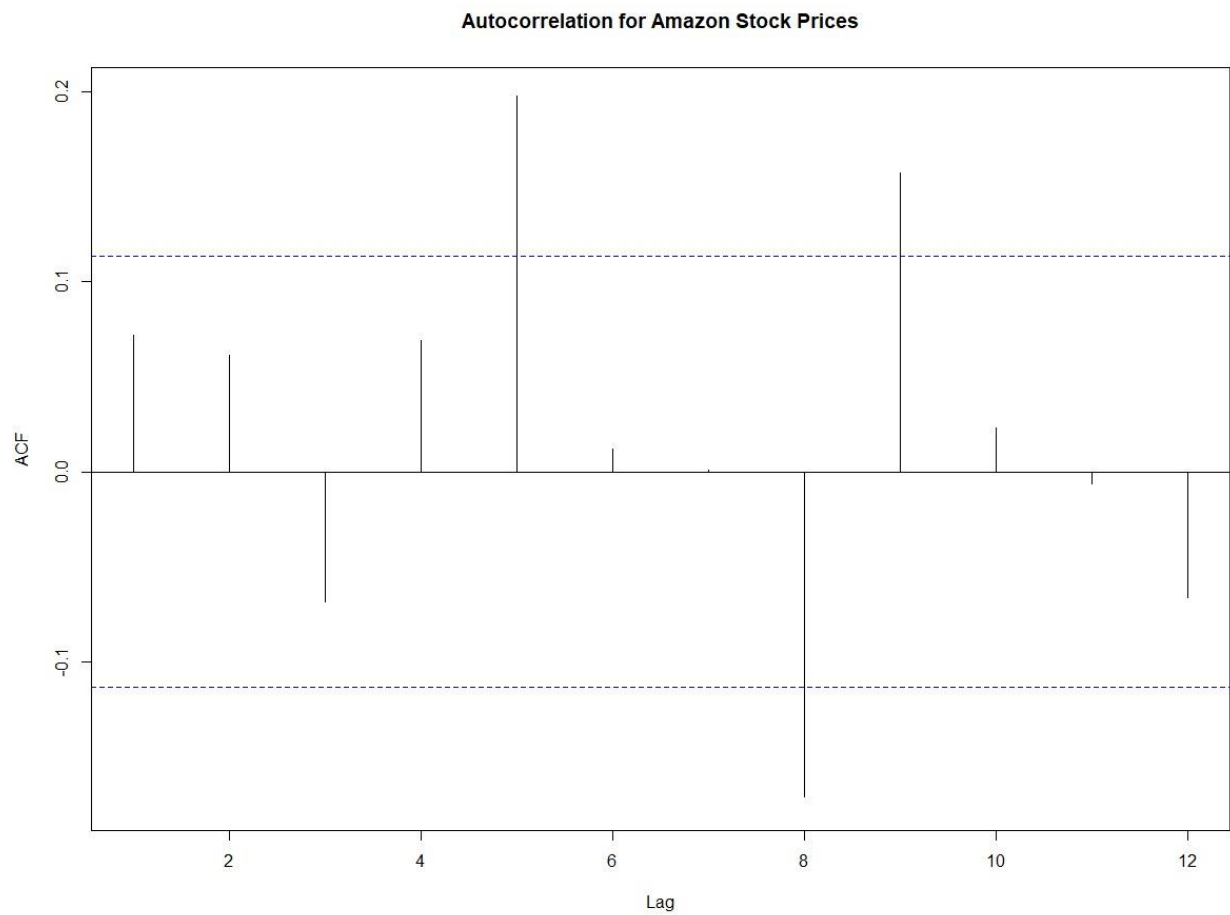

Step 1: Define the Goal - The goal of the project is to forecast the stock price of Amazon using historical data and various statistical techniques. The forecasting models developed for this project were done in R language.

Step 2: Get the Data - We have obtained the data from the Kaggle, and it deals with historical data on Amazon's stock price, containing information on the stock performance on a monthly basis has been obtained. We have data such as monthly high, monthly low, average stock price for the month. The data spans from May 1997 to the September 2023

The first column is the 'date' column, which provides the date on which the stock price was recorded. The second column is the 'price' column, which displays the opening price of the stock on the given date. The third column is the 'high' column, which represents the highest price at which the stock was traded during the day. The fourth column is the 'low' column, which displays the lowest price at which the stock was traded during the day.

Using the day high and day low, we calculated day average $((\text{day high} + \text{day low})/2)$. We then calculated monthly average by taking average of the day averages we have. Also we set the date as 15th of the month, so that we can forecast on monthly basis going forward.

Step 3: Explore and visualize the series - The project involves exploring and visualizing the historical data to identify various time series components of trends, patterns, and outliers.



Step 4: Data pre-processing - The historical data may require pre-processing steps such as handling missing data, smoothing, and scaling.

Step 5: Partition Series - The dataset may be partitioned into training and test sets for modelling and evaluating the forecast accuracy. In this model, the partition set was 80% (240 months) for training data & 20% (60 months) for validation data.

Step 6 & 7: Apply forecasting and compare performances - The project involves applying various forecasting models such as regression analysis, time series analysis, and predictive modelling. The performance of each model can be compared based on metrics such as RMSE and MAPE.

Introductory R code to start forecasting in R environment.

```
library(forecast)
library(zoo)
```

```
# See the first and last 6 records of the dataset.
head(Amazon.Data)
tail(Amazon.Data)
```

```
12
13 # See the first and last 6 records of the file for Amazon data.
14 head(Amazon.Data)
15 tail(Amazon.Data)
16
17
72:1 (Top Level) ↕
```

```
Console Terminal Background Jobs
R 4.2.1 - C:/Users/o/Desktop/My subjects/Third Semester/Capstone/Dataset/
300 15-09-2023 138.2280
> # See the first and last 6 records of the file for Amazon data.
> head(Amazon.Data)
      Month AverageStockPrice
1 15-10-1998           0.883
2 15-11-1998           1.297
3 15-12-1998           2.167
4 15-01-1999           3.367
5 15-02-1999           2.735
6 15-03-1999           3.363
> tail(Amazon.Data)
      Month AverageStockPrice
295 15-04-2023          103.345
296 15-05-2023          111.817
297 15-06-2023          126.427
298 15-07-2023          130.937
299 15-08-2023          135.543
300 15-09-2023          138.228
```

##ts() function is used to create time series data set.


```
#Amazon.ts <- ts(Amazon.Data$Average.Stock.Price,start = c(1998,10), end =
c(2023,9),freq=12)
```

```
> Amazon.ts <- ts(Amazon.Data$AverageStockPrice,
+               start = c(1998, 10), end = c(2023,9) , freq = 12)
> Amazon.ts
```

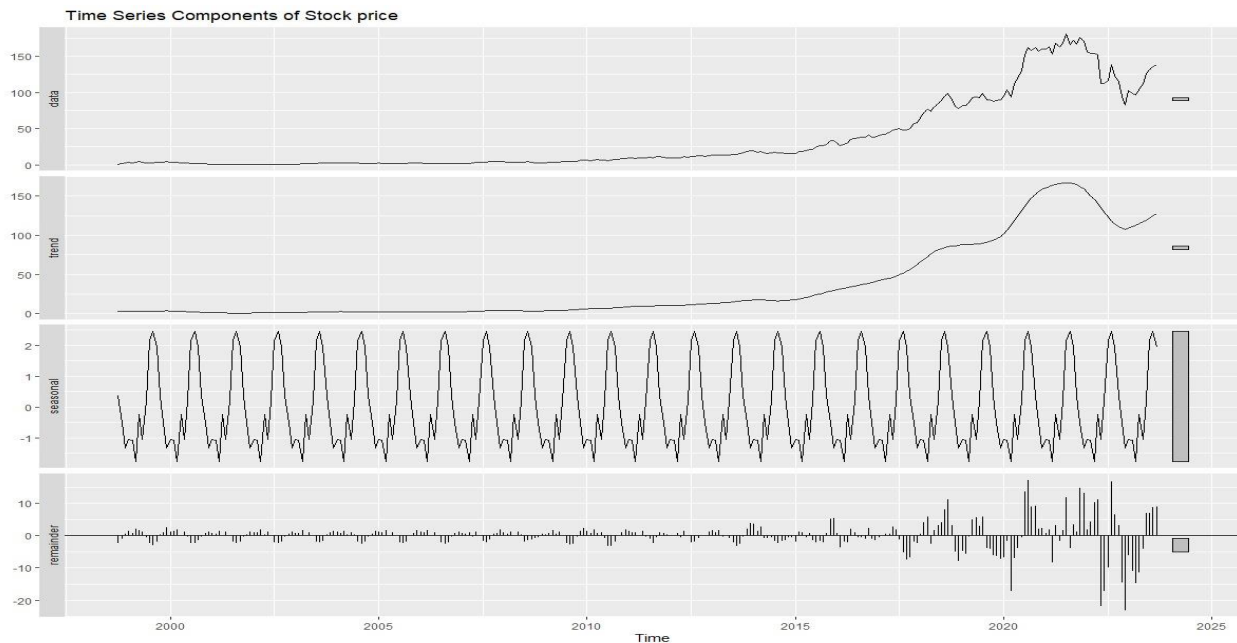
	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1998										0.8830	1.2970	2.1670
1999	3.3670	2.7350	3.3630	4.5530	3.3610	2.7860	3.0010	2.6400	3.2460	3.9800	3.8100	4.6080
2000	3.4190	3.6190	3.3200	2.8250	2.6720	2.2800	1.8370	1.8050	2.0970	1.5310	1.5210	1.0570
2001	0.8900	0.6810	0.5490	0.6700	0.7880	0.7120	0.7390	0.5250	0.3720	0.3810	0.4360	0.5580
2002	0.6010	0.6530	0.7660	0.7250	0.9060	0.8660	0.7390	0.7230	0.8140	0.9160	1.0650	1.0940
2003	1.0700	1.0700	1.2460	1.3130	1.5960	1.7640	1.9350	2.1000	2.3670	2.7950	2.6330	2.5650
2004	2.6930	2.2730	2.1200	2.3250	2.1610	2.5360	2.3100	1.8960	2.0290	1.9180	1.9140	2.0300
2005	2.1220	1.8080	1.7260	1.6840	1.7290	1.7340	1.7590	1.7560	1.8740	2.1820	2.2280	2.4400
2006	2.2610	1.9380	1.8220	1.8170	1.7050	1.7590	1.6700	1.3850	1.5740	1.6930	2.0170	1.9690
2007	1.8790	1.9710	1.9300	2.3110	3.2180	3.5180	3.7230	3.8400	4.4030	4.5700	4.1510	4.5790
2008	4.0860	3.6300	3.4220	3.8450	3.8680	4.0000	3.6160	4.1080	3.8150	2.8300	2.2270	2.4660
2009	2.6070	3.1600	3.4210	3.9230	3.8850	4.1600	4.1700	4.2000	4.3250	5.0550	6.4030	6.7610
2010	6.3860	5.8890	6.5240	7.0520	6.3810	6.0980	5.8380	6.3000	7.3490	8.0210	8.4180	8.9460
2011	9.1105	9.0470	8.4350	9.2000	9.9050	9.5820	10.7490	9.9670	11.1370	11.3170	10.2880	9.2030
2012	9.2250	9.1280	9.4540	9.7580	11.0610	10.9440	11.1870	11.9470	12.7850	12.2420	11.7020	12.6380
2013	13.4330	13.1930	13.2600	13.1560	13.1270	13.6970	14.9160	14.5670	15.2130	16.2610	18.2090	19.5730
2014	19.7600	17.6830	18.1760	16.0620	15.1240	16.2250	16.9570	16.3670	16.6060	15.4100	15.8930	15.4730
2015	15.1450	18.7420	18.7800	19.7270	21.3630	21.6300	23.9790	25.9700	26.0470	28.2300	32.8360	33.4590
2016	30.0780	26.5250	28.5170	30.6690	34.8280	35.7970	37.0550	38.2120	38.3890	41.2650	38.2090	38.1760
2017	40.2550	41.7160	42.6720	45.2250	48.0380	49.5150	50.3940	48.5780	48.4760	50.0090	56.8970	58.4640
2018	65.2190	72.1620	76.8450	73.5120	79.6010	84.9070	89.1620	94.7400	98.2700	89.4120	81.1700	78.2450
2019	81.7500	81.3060	86.0660	93.1330	93.4730	92.6200	98.1500	89.7200	89.9250	87.5070	88.7130	89.2740
2020	94.2090	103.4830	93.3880	111.0700	119.5840	130.5250	152.4850	162.3110	158.2720	161.7300	157.0490	159.8240
2021	160.3220	163.2500	153.5360	167.5910	162.6750	168.1900	180.5600	165.4550	171.4710	166.6350	176.2980	170.9810
2022	155.6440	153.7990	154.2340	151.7980	112.3100	112.5740	116.4810	137.8070	123.1970	114.3210	94.2520	83.2600
2023	102.2400	99.0360	96.3980	103.3450	111.8170	126.4270	130.9370	135.5430	138.2280			

#stl() component is used to plot time series components(seasonality, trend and level) of the original data:

##(Visualization)

```
Stock.stl<- stl(Amazon.ts, s.window="periodic")
```

```
autoplot(Stock.stl, main="Time Series Components of Stock price")
```



```
# Use Arima() function to fit AR(1) model.
```

```
# The ARIMA model of order = c(1,0,0) gives an AR(1) model.
Stock.ar1<- Arima(Amazon.ts, order = c(1,0,0))
summary(Stock.ar1)
```

```
33 # Use Arima() function to fit AR(1) model.
34 # The ARIMA model of order = c(1,0,0) gives an AR(1) model.
35 # ARIMA (1,0,0) model is an autoregressive model of order 1 without differencing and without average moving component.
36 Stock.ar1<- Arima(Amazon.ts, order = c(1,0,0))
37 summary(Stock.ar1)
```

```
R 4.2.1 - C:/Users/o/Desktop/My subjects/Third Semester/Capstone/Dataset/
> # The ARIMA model of order = c(1,0,0) gives an AR(1) model.
> # ARIMA (1,0,0) model is an autoregressive model of order 1 without differencing and without average moving component.
> Stock.ar1<- Arima(Amazon.ts, order = c(1,0,0))
> summary(Stock.ar1)
Series: Amazon.ts
ARIMA(1,0,0) with non-zero mean

Coefficients:
      ar1      mean
    0.9971  58.3698
s.e.  0.0032  53.1644

sigma^2 = 23.51: log likelihood = -900.85
AIC=1807.7  AICc=1807.78  BIC=1818.81

Training set error measures:
      ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set 0.3742477 4.832293 2.191709 -5.561658 11.73532 0.2181803 0.07629686
```

The above code `Stock.ar1<- Arima(Amazon.ts, order = c(1,0,0))` fits an ARIMA (Autoregressive Integrated Moving Average) model of order (1,0,0) to the Amazon.ts time series data. This model is an AR(1) model, which means it has a single autoregressive term and no moving average term.

The `summary(Stock.ar1)` command provides a summary of the model's coefficients and other key statistics.

The output shows that the AR(1) model has an autoregressive coefficient (AR1) of 0.9971 and a mean of 58.3698. The standard error of the coefficient estimate (s.e.) is 0.0032, and the sigma squared is estimated to be 23.51. The log-likelihood of the model is -900.85, and the Akaike information criterion (AIC), corrected AIC (AICc), and Bayesian information criterion (BIC) values are also provided.

###Checking predictability using hypothesis testing method:

#Predictability test approach-1

```
# Apply z-test to test the null hypothesis that beta
# coefficient of AR(1) is equal to 1.
```

```

ar1 <- 0.9971
s.e.<- 0.0032
null_mean<- 1
alpha <- 0.05
z.stat<- (ar1-null_mean)/s.e.
z.stat
p.value<- pnorm(z.stat)
p.value
if (p.value<alpha) {
  "Reject null hypothesis"
} else {
  "Accept null hypothesis"
}

```

```

42
43 # Apply z-test to test the null hypothesis that beta
44 # coefficient of AR(1) is equal to 1.
45 ar1 <- 0.9971
46 s.e. <- 0.0032
47 null_mean <- 1
48 alpha <- 0.05
49 z.stat <- (ar1-null_mean)/s.e.
50 z.stat
51 p.value <- pnorm(z.stat)
52 p.value
53 if (p.value<alpha) {
54   "Reject null hypothesis"
55 } else {
56   "Accept null hypothesis"
57 }
58

```

59:1 (Top Level) ↕

Console Terminal × Background Jobs ×

R 4.2.1 · C:/Users/o/Desktop/My subjects/Third Semester/Capstone/Dataset/ ↗

```

Training set 0.3742477 4.832293 2.191709 -5.561658 11.73532 0.2181803 0.07629686
> # Apply z-test to test the null hypothesis that beta
> # coefficient of AR(1) is equal to 1.
> ar1 <- 0.9971
> s.e. <- 0.0032
> null_mean <- 1
> alpha <- 0.05
> z.stat <- (ar1-null_mean)/s.e.
> z.stat
[1] -0.90625
> p.value <- pnorm(z.stat)
> p.value
[1] 0.1824018
> if (p.value<alpha) {
+   "Reject null hypothesis"
+ } else {
+   "Accept null hypothesis"
+ }
[1] "Accept null hypothesis"

```

This code is performing a predictability test on the AR(1) model for Amazon stock price. The null hypothesis is that the beta coefficient of AR(1) is equal to 1, which means that the stock price is predictable based on its past values. The alternative hypothesis is that stock price is not predictable based on its past values and it is a random walk.

The code first sets the value of the AR(1) coefficient (ar1) to 0.9971 and the standard error (s.e.) to 0.0032. The null hypothesis is set to a value of 1. The significance level (alpha) is set to 0.05.

The z-statistic is calculated as $(ar1 - \text{null_mean}) / \text{s.e.}$, which gives a value of -0.90625. The p-value is then calculated using the `pnorm()` function, which gives a value of 0.1824.

The code then checks if the p-value is greater than the significance level (alpha). **Since the p-value is greater than the significance level, the null hypothesis is accepted, which means that the stock price is random and not predictable.**

#Predictability test approach-2

Create first difference of ClosePrice data using diff() function.

diff.stock.price<- diff(Amazon.ts, lag = 1)

diff.stock.price

```

58 #Predictability test approach-2
59
60 # Create first difference of ClosePrice data using diff() function.
61 diff.stock.price <- diff(Amazon.ts, lag = 1)
62 diff.stock.price
63
64

```

65:1 (Top Level) ±

Console Terminal Background Jobs

R 4.2.1 - C:/Users/o/Desktop/My subjects/Third Semester/Capstone/Dataset/

```

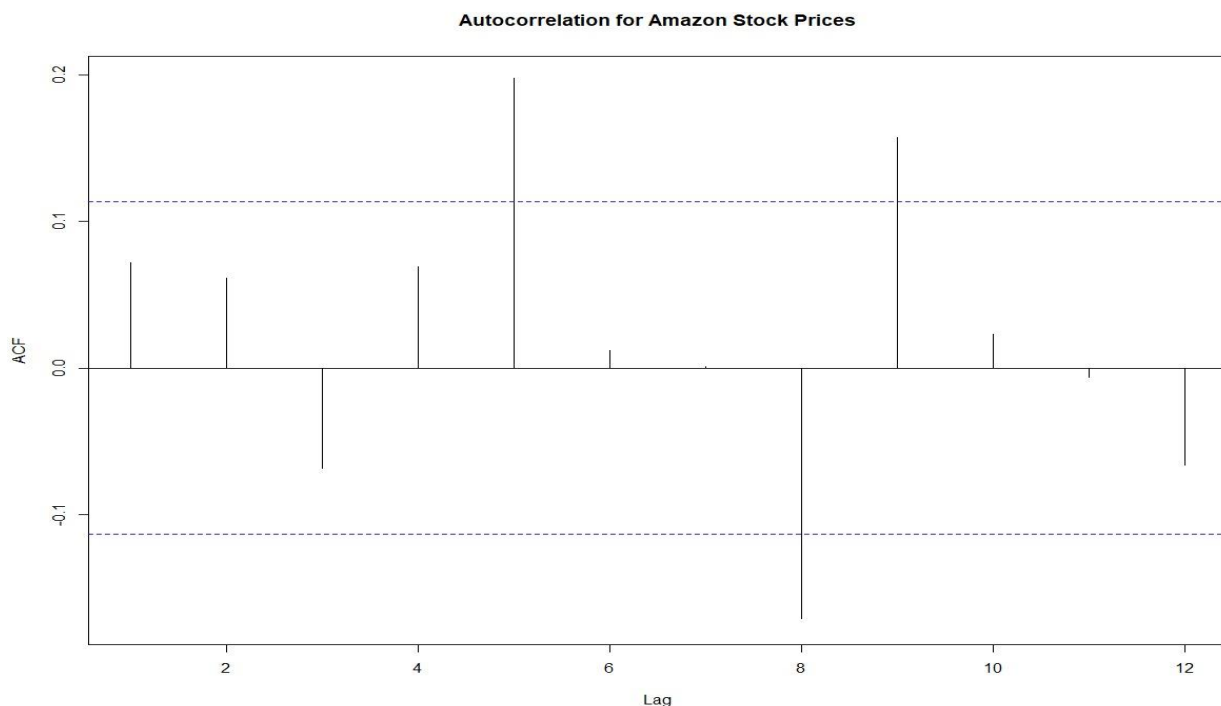
+ }
[1] "Accept null hypothesis"
> # Create first difference of ClosePrice data using diff() function.
> diff.stock.price <- diff(Amazon.ts, lag = 1)
> diff.stock.price

```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1998											0.4140	0.8700
1999	1.2000	-0.6320	0.6280	1.1900	-1.1920	-0.5750	0.2150	-0.3610	0.6060	0.7340	-0.1700	0.7980
2000	-1.1890	0.2000	-0.2990	-0.4950	-0.1530	-0.3920	-0.4430	-0.0320	0.2920	-0.5660	-0.0100	-0.4640
2001	-0.1670	-0.2090	-0.1320	0.1210	0.1180	-0.0760	0.0270	-0.2140	-0.1530	0.0090	0.0550	0.1220
2002	0.0430	0.0520	0.1130	-0.0410	0.1810	-0.0400	-0.1270	-0.0160	0.0910	0.1020	0.1490	0.0290
2003	-0.0240	0.0000	0.1760	0.0670	0.2830	0.1680	0.1710	0.1650	0.2670	0.4280	-0.1620	-0.0680
2004	0.1280	-0.4200	-0.1530	0.2050	-0.1640	0.3750	-0.2260	-0.4140	0.1330	-0.1110	-0.0040	0.1160
2005	0.0920	-0.3140	-0.0820	-0.0420	0.0450	0.0050	0.0250	-0.0030	0.1180	0.3080	0.0460	0.2120
2006	-0.1790	-0.3230	-0.1160	-0.0050	-0.1120	0.0540	-0.0890	-0.2850	0.1890	0.1190	0.3240	-0.0480
2007	-0.0900	0.0920	-0.0410	0.3810	0.9070	0.3000	0.2050	0.1170	0.5630	0.1670	-0.4190	0.4280
2008	-0.4930	-0.4560	-0.2080	0.4230	0.0230	0.1320	-0.3840	0.4920	-0.2930	-0.9850	-0.6030	0.2390
2009	0.1410	0.5530	0.2610	0.5020	-0.0380	0.2750	0.0100	0.0300	0.1250	0.7300	1.3480	0.3580
2010	-0.3750	-0.4970	0.6350	0.5280	-0.6710	-0.2830	-0.2600	0.4620	1.0490	0.6720	0.3970	0.5280
2011	0.1645	-0.0635	-0.6120	0.7650	0.7050	-0.3230	1.1670	-0.7820	1.1700	0.1800	-1.0290	-1.0850
2012	0.0220	-0.0970	0.3260	0.3040	1.3030	-0.1170	0.2430	0.7600	0.8380	-0.5430	-0.5400	0.9360
2013	0.7950	-0.2400	0.0670	-0.1040	-0.0290	0.5700	1.2190	-0.3490	0.6460	1.0480	1.9480	1.3640
2014	0.1870	-2.0770	0.4930	-2.1140	-0.9380	1.1010	0.7320	-0.5900	0.2390	-1.1960	0.4830	-0.4200
2015	-0.3280	3.5970	0.0380	0.9470	1.6360	0.2670	2.3490	1.9910	0.0770	2.1830	4.6060	0.6230
2016	-3.3810	-3.5530	1.9920	2.1520	4.1590	0.9690	1.2580	1.1570	0.1770	2.8760	-3.0560	-0.0330
2017	2.0790	1.4610	0.9560	2.5530	2.8130	1.4770	0.8790	-1.8160	-0.1020	1.5330	6.8880	1.5670
2018	6.7550	6.9430	4.6830	-3.3330	6.0890	5.3060	4.2550	5.5780	3.5300	-8.8580	-8.2420	-2.9250
2019	3.5050	-0.4440	4.7600	7.0670	0.3400	-0.8530	5.5300	-8.4300	0.2050	-2.4180	1.2060	0.5610
2020	4.9350	9.2740	-10.0950	17.6820	8.5140	10.9410	21.9600	9.8260	-4.0390	3.4580	-4.6810	2.7750
2021	0.4980	2.9280	-9.7140	14.0550	-4.9160	5.5150	12.3700	-15.1050	6.0160	-4.8360	9.6630	-5.3170
2022	-15.3370	-1.8450	0.4350	-2.4360	-39.4880	0.2640	3.9070	21.3260	-14.6100	-8.8760	-20.0690	-10.9920
2023	18.9800	-3.2040	-2.6380	6.9470	8.4720	14.6100	4.5100	4.6060	2.6850			

In this code, the first difference of the Amazon.ts time series data is calculated using the diff() function. The lag argument is set to 1, which means that the first difference will be taken with respect to the previous value. The resulting data is printed to the console, showing the differences in the data from one time period to the next. This approach can be used to test for predictability in the time series data, as it can reveal whether the changes in the data are random or follow some pattern or trend.

```
#UseAcf()functionto identify autocorrelation for first differenced Stock Price and plot autocorrelation for different lags  
# (up to maximum of 12).  
Acf(diff.stock.price, lag.max = 12,  
main = "Autocorrelation for Amazon Stock Prices")
```



We can see significant correlations existing at lags 5, 8&9. Therefore, we can say that this is not random walk and we can predict the data.

Creating Time Series Partition.

```
# Define the numbers of months in the training and validation sets,  
# nTrain and nValid, respectively.  
# Total number of period length(ridership.ts) = 300.  
# nTrain = 240 months, from October 1998 to September 2018.  
# nvalid = 60 months, from October 2018 to September 2023.
```



```

nValid<- 60
nTrain<- length(Amazon.ts) - nValid
train.ts<- window(Amazon.ts, start = c(1998, 10), end = c(1998, nTrain))
valid.ts<- window(Amazon.ts, start = c(1998, nTrain + 1),
  end = c(1998, nTrain + nValid))

```

MODEL ONE:

Use two-level model, regression with linear trend and Seasonality and trailing ma for residuals for entire data set.

Use two-level (combined) forecast to forecast 12 future periods from September 2023.

Fit a regression model with linear trend and seasonality for entire data set.

```

tot.trend.seas<- tslm(Amazon.ts ~ trend + season)
summary(tot.trend.seas)

```

```

87 #####MODEL ONE:
88
89
90 ## USE TWO-LEVEL MODEL, REGRESSION WITH LINEAR TREND AND
91 ## SEASONALITY AND TRAILING MA FOR RESIDUALS FOR ENTIRE
92 ## DATA SET. USE TWO-LEVEL (COMBINED) FORECAST TO FORECAST
93 ## 12 FUTURE PERIODS FROM APRIL 2023.
94
95 # Fit a regression model with linear trend and seasonality for
96 # entire data set.
97 tot.trend.seas <- tslm(Amazon.ts ~ trend + season)
98 summary(tot.trend.seas)
98:24 (Top Level)

```

Console **Terminal** **Background Jobs**

```

R 4.2.1 · C:/Users/o/Desktop/My subjects/Third Semester/Capstone/Dataset/
> # Fit a regression model with linear trend and seasonality for
> # entire data set.
> tot.trend.seas <- tslm(Amazon.ts ~ trend + season)
> summary(tot.trend.seas)

```

Call:
tslm(formula = Amazon.ts ~ trend + season)

Residuals:

	Min	1Q	Median	3Q	Max
	-40.338	-24.717	-3.177	16.730	86.056

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-35.86983	6.68291	-5.367	1.65e-07 ***
trend	0.46609	0.01999	23.315	< 2e-16 ***
season2	-0.03707	8.47446	-0.004	0.997
season3	-0.84423	8.47453	-0.100	0.921
season4	0.58244	8.47465	0.069	0.945
season5	-0.36816	8.47481	-0.043	0.965
season6	0.75255	8.47502	0.089	0.929
season7	2.66659	8.47528	0.315	0.753
season8	2.92502	8.47559	0.345	0.730
season9	2.45169	8.47594	0.289	0.773
season10	1.53116	8.47465	0.181	0.857
season11	0.60675	8.47453	0.072	0.943
season12	-0.25237	8.47446	-0.030	0.976

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 29.96 on 287 degrees of freedom
Multiple R-squared: 0.6553, Adjusted R-squared: 0.6409
F-statistic: 45.47 on 12 and 287 DF, p-value: < 2.2e-16

The output is the summary of a multiple linear regression model using the `tslm()` function from the forecast package. The response variable is the Amazon stock price (`Amazon.ts`), and the predictors are a linear trend (`trend`) and seasonal dummies (`season2`, `season3`, ..., `season12`), which represent the 12 months of the year. The estimates for the coefficients are given in the "Coefficients" table.

The intercept estimate is -35.869, which represents the expected value of the response variable when all predictors are equal to zero. The estimate for the trend variable is 0.466, which represents the expected change in the response variable for a unit increase in time (in months). The estimates for the seasonal dummies represent the expected difference in the response variable between the given month and the base month (January, represented by `season1`).

The p-values in the "Pr(>|t|)" column indicate the significance of each coefficient, with smaller p-values indicating stronger evidence against the null hypothesis (that the coefficient is zero). The p-value for trend is very small, indicating that it is highly significant. However, most of the seasonal dummies have large p-values, indicating that they are not significant predictors of the response variable.

Create regression forecast for future 12 periods.

`tot.trend.seas.pred <- forecast(tot.trend.seas, h = 12, level = 0)`

`tot.trend.seas.pred`

```
> # Create regression forecast for future 12 periods.
> tot.trend.seas.pred <- forecast(tot.trend.seas, h = 12,
+                               level = 0)
> tot.trend.seas.pred
```

	Point Forecast	Lo 0	Hi 0
Oct 2023	105.9531	105.9531	105.9531
Nov 2023	105.4948	105.4948	105.4948
Dec 2023	105.1018	105.1018	105.1018
Jan 2024	105.8202	105.8202	105.8202
Feb 2024	106.2493	106.2493	106.2493
Mar 2024	105.9082	105.9082	105.9082
Apr 2024	107.8009	107.8009	107.8009
May 2024	107.3164	107.3164	107.3164
Jun 2024	108.9032	108.9032	108.9032
Jul 2024	111.2833	111.2833	111.2833
Aug 2024	112.0079	112.0079	112.0079
Sep 2024	112.0006	112.0006	112.0006

The above output shows the Point Forecast column of the predicted values for each period, while the Lo 0 and Hi 0 columns show the lower and upper bounds of the 0% prediction interval, respectively. The forecast predicts an overall upward trend, with some fluctuations around the trend. It is important to note that this forecast assumes that the underlying trend and seasonality patterns will continue into the future, and that there will be no major unexpected events that could affect the stock price.

Identify and display regression residuals for entire data set.

```
tot.trend.seas.res <- tot.trend.seas$residuals
tot.trend.seas.res
```

```
> # Identify and display regression residuals for entire data set.
> tot.trend.seas.res <- tot.trend.seas$residuals
> tot.trend.seas.res
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1998										34.7555838	35.6279038	36.8909438
1999	37.3724838	36.3114638	37.2805438	36.5777838	35.8703038	33.7085038	31.5433838	30.4578638	31.0711038	32.2595552	32.5478752	33.7389152
2000	31.8314552	31.6024352	31.6445152	29.2567552	29.5882752	27.6094752	24.7863552	24.0298352	24.3290752	24.2175265	24.6658465	24.5948865
2001	23.7094265	23.0714065	23.2804865	21.5087265	22.1112465	20.4484465	18.0953265	17.1568065	17.0110465	17.4744979	17.9878179	18.5028579
2002	17.8273979	17.4503779	17.9044579	15.9706979	16.6362179	15.0094179	12.5022979	11.7617779	11.8600179	12.4164692	13.0237892	13.4458292
2003	12.7033692	12.2743492	12.7914292	10.9656692	11.7331892	10.3143892	8.1052692	7.5457492	7.8199892	8.7024406	8.9987606	9.3238006
2004	8.7333406	7.8843206	8.0724006	6.3846406	6.7051606	5.4933606	2.8872406	1.7487206	1.8889606	2.2324119	2.6867319	3.1957719
2005	2.5693119	1.8262919	2.0853719	0.1506119	0.6801319	-0.9016681	-3.2567881	-3.9843081	-3.8590681	-3.0966167	-2.5922967	-1.9872567
2006	-2.8847167	-3.6367367	-3.4116567	-5.3094167	-4.9368967	-6.4696967	-8.9388167	-9.9483367	-9.7520967	-9.1786454	-8.3963254	-8.0512854
2007	-8.8597454	-9.1967654	-8.8966854	-10.4084454	-9.0169254	-10.3037254	-12.4788454	-13.0863654	-12.5161254	-11.8946740	-11.8553540	-11.0343140
2008	-12.2457740	-13.1307940	-12.9977140	-14.4674740	-13.9599540	-15.4147540	-18.1788740	-18.4113940	-18.6971540	-19.2277027	-19.3723827	-18.7403427
2009	-19.3178027	-19.1938227	-18.5917427	-19.9825027	-19.5359827	-20.8477827	-23.2179027	-23.9124227	-23.7801827	-22.5957313	-20.7894113	-20.0383713
2010	-21.1318313	-22.0578513	-21.0817713	-22.4465313	-22.6330113	-24.5028113	-27.1429313	-27.4054513	-26.3492113	-25.2227600	-24.3674400	-23.4464000
2011	-24.0003600	-24.4928800	-24.7638000	-25.8915600	-24.7020400	-26.6118400	-27.8249600	-29.3314800	-28.1542400	-27.5197887	-28.0904687	-28.7824287
2012	-29.4788887	-30.0049087	-29.3378287	-30.9265887	-29.1390687	-30.8428687	-32.9799887	-32.9445087	-32.0992687	-32.1878173	-32.2694973	-30.9404573
2013	-30.8639173	-31.5329373	-31.1248573	-33.1216173	-32.6660973	-33.6828973	-34.8440173	-35.9175373	-35.2642973	-33.7618460	-31.3555260	-29.5984860
2014	-30.1299460	-32.6359660	-31.8018860	-35.8086460	-36.2621260	-36.7479260	-38.3960460	-39.7105660	-39.4643260	-40.2058746	-39.2645546	-39.2915146
2015	-40.3379746	-37.1699946	-36.7909146	-37.7366746	-35.6161546	-36.9359546	-36.9670746	-35.7005946	-35.6163546	-32.9789033	-27.9145833	-26.8985433
2016	-30.9980033	-34.9800233	-32.6469433	-32.3877033	-27.7441833	-28.3619833	-29.4841033	-29.0516233	-28.8673833	-25.5369319	-28.1346119	-27.7745719
2017	-26.4140319	-25.3820519	-24.0849719	-23.4247319	-20.1272119	-20.2370119	-21.7381319	-24.2786519	-24.3734119	-22.3859606	-15.0396406	-13.0796006
2018	-7.0430606	-0.5290806	4.4949994	-0.7307606	5.8427594	9.5619594	11.4368394	16.2903194	19.8275594	11.4240108	3.6403308	1.1083708
2019	3.8949108	3.0218908	8.1229708	13.2972108	14.1217308	11.6819308	14.8318108	5.6772908	5.8895308	3.9259821	5.5903021	6.5443421
2020	10.7608821	19.6058621	9.8519421	25.6411821	34.6397021	43.9939021	63.5737821	72.6752621	68.6435021	72.5559535	68.3332735	71.5013135
2021	71.2808535	73.7798335	64.4069135	76.5691535	72.1376735	76.0658735	86.0557535	70.2262335	76.2494735	71.8679248	81.9892448	77.0652848
2022	61.0098248	58.7358048	59.5118848	55.1831248	16.1796448	14.8568448	16.3837248	36.9852048	22.3824448	13.9608962	-5.6497838	-16.2487438
2023	2.0127962	-1.6202238	-3.9171438	1.1370962	10.0936162	23.1168162	25.2466962	29.1281762	31.8204162			

These residuals are the differences between the actual values of the time series and the values predicted by the trend and seasonal components of the model. They represent the unexplained variability in the data after accounting for the trend and seasonal patterns. The residuals are displayed in the output above, which shows the values for each month from October 1998 to September 2023. The blank spaces in the last row indicate that the time series model has not yet made predictions for those months.

Use trailing MA to forecast residuals for entire data set.


```
tot.ma.trail.res <- rollmean(tot.trend.seas.res, k = 4, align = "right")
```

```
tot.ma.trail.res
```

```

100 # Use trailing MA to forecast residuals for entire data set.
101 tot.ma.trail.res <- rollmean(tot.trend.seas.res, k = 4,
102                             align = "right")
103 tot.ma.trail.res
104
105

```

```

R 4.2.1 • C:/users/manu7/Desktop/MSBA Sem-2/Sem-2/BAN 673/Final Project/OurProject/OurProject/
> # Use trailing MA to forecast residuals for entire data set.
> tot.ma.trail.res <- rollmean(tot.trend.seas.res, k = 4,
+                             align = "right")
> tot.ma.trail.res

```

	Jan	Feb	Mar	Apr	May	Jun	Jul
2013							11.903409091
2014	12.032159091	12.493159091	13.042409091	11.854511364	10.440113636	9.276465909	7.676818182
2015	2.023068182	0.681568182	0.433318182	-0.564579545	-1.246477273	-1.905125000	-2.852272727
2016	1.258977273	0.592477273	0.009227273	-1.408670455	-2.620568182	-3.496715909	-5.153863636
2017	-6.075113636	-6.926613636	-6.502363636	-6.990261364	-6.429659091	-5.513306818	-5.732954545
2018	-4.261704545	-4.205704545	-4.006454545	-5.416852273	-7.298750000	-7.394897727	-7.674545455
2019	-14.000795455	-15.157295455	-14.095545455	-13.125943182	-13.242840909	-14.496488636	-16.471136364
2020	-12.359886364	-10.608886364	-12.552136364	-15.187534091	-16.941931818	-18.138079545	-15.575227273
2021	-3.833977273	1.467022727	4.866272727	9.918375000	15.193977273	18.907829545	23.795681818
2022	41.726931818	39.912931818	36.479681818	31.126784091	24.232386364	17.343738636	10.084090909
2023	-16.509659091	-18.248659091	-17.674409091				

	Aug	Sep	Oct	Nov	Dec
2013	11.148159091	10.134659091	9.731909091	9.856909091	10.999659091
2014	6.839068182	6.393068182	5.515318182	4.332818182	3.513068182
2015	-2.660022727	-2.471022727	-1.508772727	-0.426272727	0.701477273
2016	-5.911613636	-6.037613636	-5.627863636	-5.740363636	-5.885113636
2017	-6.063204545	-7.324204545	-7.774454545	-7.421954545	-6.119204545
2018	-6.467295455	-6.138295455	-7.231045455	-9.606045455	-12.130795455
2019	-18.333886364	-18.599886364	-17.515136364	-16.317636364	-14.497386364
2020	-12.947977273	-12.338977273	-11.711727273	-9.501727273	-6.851477273
2021	28.640431818	34.056931818	38.011681818	41.976681818	43.549431818
2022	5.756340909	2.325340909	-1.889909091	-7.152409091	-13.279659091
2023					

Environment	History	Connections	Tutorial
<div> <div>Import Dataset</div> <div>233 MiB</div> </div>			
<div> <div>R</div> <div>Global Environment</div> </div>			
alpha	0.05		
ar1	0.9943		
diff.stock.price	Time-Series [1:119] from 2013 to 2023: 2.14 0.04 0.59 -0.62 0.2 ...		
Google.ts	Time-Series [1:120] from 2013 to 2023: 19.8 21.9 22 22.6 21.9 ...		
nTrain	96		
null_mean	1		
nvalid	24		
p.value	0.186565191332239		
s.e.	0.0064		
tot.ma.trail.res	Time-Series [1:117] from 2014 to 2023: 11.9 11.15 10.13 9.73 9.86 ...		
tot.trend.seas.res	Time-Series [1:120] from 2013 to 2023: 10.95 13.4 12.48 10.79 7.93 ...		
train.fst.2level	Time-Series [1:117] from 2014 to 2023: 23.7 25.2 22.9 22.3 24.4 ...		
train.ts	Time-Series [1:93] from 2013 to 2021: 19.8 21.9 22 22.6 21.9 ...		
valid.ts	Time-Series [1:24] from 2021 to 2023: 90 103 102 114 115 ...		
z.stat	-0.890625000000006		

This code uses a rolling mean to forecast the residuals for the entire data set. The `rollmean()` function calculates the mean of a rolling window, with a window size of 4 and the `align` parameter set to "right", which means that the mean is calculated using the **4 most recent residual values** to the right of the current data point.

The resulting `tot.ma.trail.res` object is a matrix that contains the forecasted residuals for each month in the data set. The matrix has rows for each year and columns for each month, with missing values for the last few months of the last year.

```
# Identify two-level forecast, regression with linear trend and
# seasonality and trailing MA for regression residuals, for
# entire data set (training period).
train.fst.2level <- tot.trend.seas.pred$fitted+tot.ma.trail.res
train.fst.2level
```

```
105 # Identify two-level forecast, regression with linear trend and
106 # seasonality and trailing MA for regression residuals, for
107 # entire data set (training period).
108 train.fst.2level <- tot.trend.seas.pred$fitted+tot.ma.trail.res
109 train.fst.2level
110
111
```

109:17 (Top Level) R Script

R 4.2.1 · C:/users/manu7/Desktop/MSBA Sem-2/Sem-2/BAN 673/Final Project/OurProject/OurProject/

```
> # Identify two-level forecast, regression with linear trend and
> # seasonality and trailing MA for regression residuals, for
> # entire data set (training period).
> train.fst.2level <- tot.trend.seas.pred$fitted+tot.ma.trail.res
> train.fst.2level
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct
2013							23.66225	25.15300	22.92050	22.32475
2014	26.98800	29.55200	29.22425	31.44994	29.72355	29.51990	30.19725	31.60550	29.94050	28.86975
2015	27.74050	28.50200	27.37675	29.79244	28.79855	29.09990	30.42975	32.86800	31.83800	32.60725
2016	37.73800	39.17450	37.71425	39.70994	38.18605	38.26990	38.88975	40.37800	39.03300	39.24975
2017	41.16550	42.41700	41.96425	44.88994	45.13855	47.01490	49.07225	50.98800	48.50800	47.86475
2018	53.74050	55.89950	55.22175	57.22494	55.03105	55.89490	57.89225	61.34550	60.45550	59.16975
2019	54.76300	55.70950	55.89425	60.27744	59.84855	59.55490	59.85725	60.24050	58.75550	59.64725
2020	67.16550	71.01950	68.19925	68.97744	66.91105	66.67490	71.51475	76.38800	75.77800	76.21225
2021	86.45300	93.85700	96.37925	104.84494	109.80855	114.48240	121.64725	128.73800	132.93550	136.69725
2022	142.77550	143.06450	138.75425	136.81494	129.60855	123.67990	118.69725	116.61550	111.96550	107.55725
2023	95.30050	95.66450	95.36175							
	Nov	Dec								
2013	24.37575	25.61350								
2014	29.61325	28.88850								
2015	35.61575	36.83850								
2016	41.06325	41.01350								
2017	50.14325	51.54100								
2018	58.72075	56.29100								
2019	62.77075	64.68600								
2020	80.34825	83.09350								
2021	142.58825	144.25600								
2022	104.22075	98.18850								
2023										

The code is generating a two-level forecast by adding the fitted values obtained from a regression model with a linear trend and seasonality to the trailing moving average of the regression residuals. The resulting forecast is then calculated for the entire dataset, including the training period. The code is stored in the object `train.fst.2level` and the resulting forecast is shown as a table of values for each month from January to December for the years 1998 to 2023.

Specifically, `tot.trend.seas.pred$fitted` is the fitted values obtained from a regression model with a linear trend and seasonality, and `tot.ma.trail.res` is the trailing moving average of the regression residuals. These two components are added together to obtain the two-level forecast for the training period. The resulting forecast is shown as a table of values for each month from January to December for the years 2013 to 2023.

Create forecast for trailing MA residuals for future 12 periods.

`tot.ma.trail.res.pred <- forecast(tot.ma.trail.res, h = 12,`

`level = 0)`

`tot.ma.trail.res.pred`

```
> # Create forecast for trailing MA residuals for future 12 periods.
> tot.ma.trail.res.pred <- forecast(tot.ma.trail.res, h = 12,
+                                 level = 0)
> tot.ma.trail.res.pred
```

	Point	Forecast	Lo 0	Hi 0
Oct 2023	31.71293	31.71293	31.71293	31.71293
Nov 2023	35.25265	35.25265	35.25265	35.25265
Dec 2023	38.11012	38.11012	38.11012	38.11012
Jan 2024	40.41682	40.41682	40.41682	40.41682
Feb 2024	42.27892	42.27892	42.27892	42.27892
Mar 2024	43.78211	43.78211	43.78211	43.78211
Apr 2024	44.99557	44.99557	44.99557	44.99557
May 2024	45.97515	45.97515	45.97515	45.97515
Jun 2024	46.76592	46.76592	46.76592	46.76592
Jul 2024	47.40427	47.40427	47.40427	47.40427
Aug 2024	47.91958	47.91958	47.91958	47.91958
Sep 2024	48.33557	48.33557	48.33557	48.33557

The forecast for the trailing MA residuals for the next 12 periods is obtained using the `forecast()` function. The resulting forecast shows the point forecast for each month as well as the lower and upper bounds of the forecast interval at level 0. The forecast values suggest that the residuals are expected to increase gradually in the next 12 months.

Develop 2-level forecast for future 12 periods by combining

regression forecast and trailing MA for residuals for future

12 periods.

`tot.fst.2level <- tot.trend.seas.pred$mean +`

`tot.ma.trail.res.pred$mean`

`tot.fst.2level`

```

> # Develop 2-level forecast for future 12 periods by combining
> # regression forecast and trailing MA for residuals for future
> # 12 periods.
> tot.fst.2level <- tot.trend.seas.pred$mean +
+   tot.ma.trail.res.pred$mean
> tot.fst.2level

```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
2023										137.6661	140.7475	143.2119
2024	146.2371	148.5282	149.6903	152.7965	153.2916	155.6691	158.6876	159.9274	160.3362			

The 2-level forecast is the sum of the point forecasts for the regression model and the trailing MA for residuals model for each period in the forecast horizon.

```

# Create a table with regression forecast, trailing MA for residuals,
# and total forecast for future 12 periods.
future12.df <- round(data.frame(tot.trend.seas.pred$mean,
tot.ma.trail.res.pred$mean, tot.fst.2level), 3)
names(future12.df) <- c("Regression.Fst", "MA.Residuals.Fst",
"Combined.Fst")
future12.df

```

Amazon.Data	300 obs. of 2 variables
future12.df	12 obs. of 3 variables
Hw.ZZZ	List of 19
Hw.ZZZ.pred	List of 10
lin.season	List of 16
lin.season.pred	List of 11
Stock.ar1	List of 18
Stock.stl	List of 8
tot.ma.trail.res.pred	List of 10
tot.trend.seas	List of 16
tot.trend.seas.pred	List of 11
values	
alpha	0.05
Amazon.ts	Time-Series [1:300] from 1999 to 2024: 0.883 1.297 2.167 3.367 2.735 ...
ar1	0.9971
diff.stock.price	Time-Series [1:299] from 1999 to 2024: 0.414 0.87 1.2 -0.632 0.628 ...
nTrain	240
null_mean	1
nvalid	60
p.value	0.182401771838493
s.e.	0.0032
tot.fst.2level	Time-Series [1:12] from 2024 to 2025: 138 141 143 146 149 ...
tot.ma.trail.res	Time-Series [1:297] from 1999 to 2024: 36.2 36.6 37 36.9 36.5 ...
tot.trend.seas.res	Time-Series [1:300] from 1999 to 2024: 34.8 35.6 36.9 37.4 36.3 ...
train.fst.2level	Time-Series [1:297] from 1999 to 2024: 2.16 2.97 3.05 4.86 4 ...
train.ts	Time-Series [1:231] from 1999 to 2018: 0.883 1.297 2.167 3.367 2.735 ...
valid.ts	Time-Series [1:60] from 2018 to 2023: 65.2 72.2 76.8 73.5 79.6 ...
z.stat	-0.906250000000004

```

> # Create a table with regression forecast, trailing MA for residuals,
> # and total forecast for future 12 periods.
> future12.df <- round(data.frame(tot.trend.seas.pred$mean,
+                                tot.ma.trail.res.pred$mean, tot.fst.2level), 3)
> names(future12.df) <- c("Regression.Fst", "MA.Residuals.Fst",
+                          "Combined.Fst")
> future12.df

```

	Regression.Fst	MA.Residuals.Fst	Combined.Fst
1	105.953	31.713	137.666
2	105.495	35.253	140.747
3	105.102	38.110	143.212
4	105.820	40.417	146.237
5	106.249	42.279	148.528
6	105.908	43.782	149.690
7	107.801	44.996	152.797
8	107.316	45.975	153.292
9	108.903	46.766	155.669
10	111.283	47.404	158.688
11	112.008	47.920	159.927
12	112.001	48.336	160.336

Now we have a table that shows the regression forecast, trailing MA for residuals forecast, and combined forecast for the future 12 periods. The combined forecast is obtained by adding the regression forecast and the trailing MA for residuals forecast. The combined forecast can be more accurate than each forecast method on its own, as it takes into account the strengths of each method.

**# Plot historical data set and two-level model's forecast for
entire data set and future 12 monthly periods.**

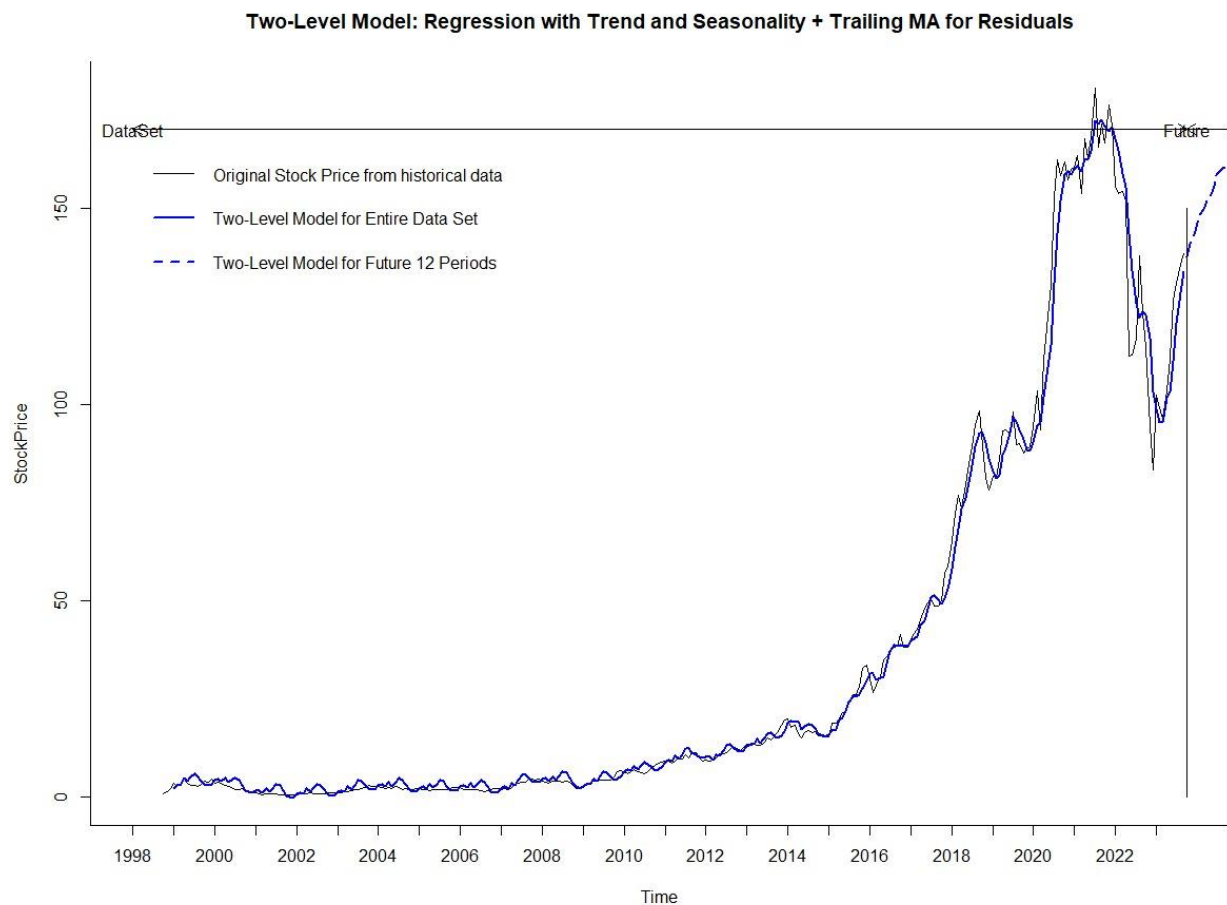
```

plot(Amazon.ts,
xlab = "Time", ylab = "StockPrice", ylim = c(0, 180),
bty = "l", xlim = c(1998, 2023.75), lwd = 1, xaxt = "n",
     main = "Two-Level Model: Regression with Trend and Seasonality + Trailing MA for
Residuals")
axis(1, at = seq(1998, 2023.75, 1), labels = format(seq(1998, 2023.75, 1)))
lines(train.fst.2level, col = "blue", lwd = 2)
lines(tot.fst.2level, col = "blue", lty = 5, lwd = 2)
legend(1998, 170, legend = c("Original Stock Price from historical data", "Two-Level
Model for Entire Data Set",
                             "Two-Level Model for Future 12 Periods"),
      col = c("black", "blue", "blue"),
lty = c(1, 1, 2), lwd = c(1, 2, 2), bty = "n")

```

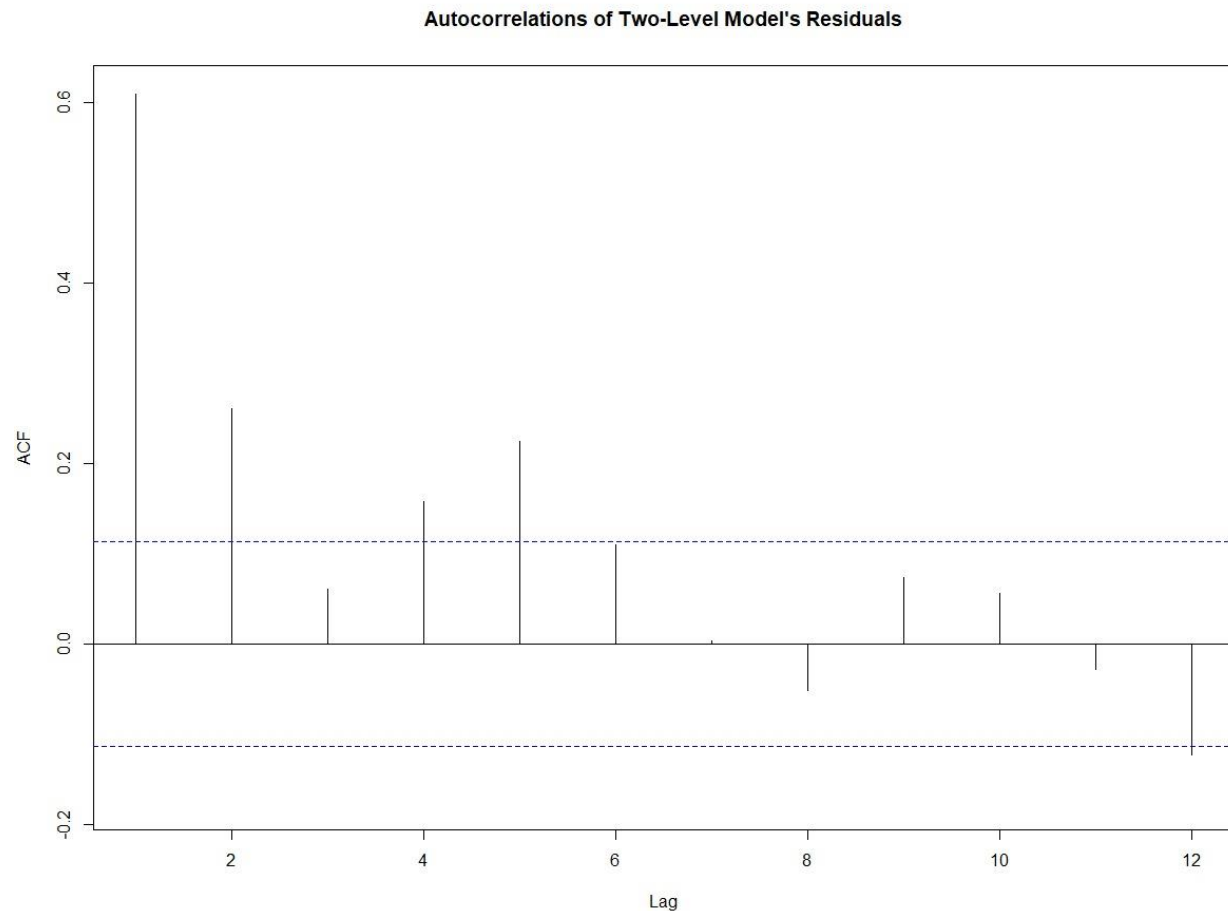
```
# Plot on chart vertical lines and horizontal arrows describing entire data set and future  
# prediction intervals.
```

```
lines(c(2023.25, 2023.25), c(0, 150))  
text(2017, 150, "DataSet")  
text(2023.5, 150, "Future")  
arrows(2013.25, 140, 2023.25, 140, code = 3, length = 0.1,  
lwd = 1, angle = 30)  
arrows(2023.25, 140, 2024, 140, code = 3, length = 0.1,  
lwd = 1, angle = 30)
```



The code is creating a plot of the historical stock price data along with the two-level model's forecast for the entire data set and the future 12 monthly periods . It also includes vertical lines and horizontal arrows to describe the entire data set and future prediction intervals, Overall, this code creates a well-labeled and informative plot of the historical data and future predictions using the two-level model.


```
# Use Acf() function to create autocorrelation chart of  
# two-level model's residuals.  
Acf((Amazon.ts - train.fst.2level), lag.max = 12,  
    main = "Autocorrelations of Two-Level Model's Residuals")
```



This will produce a plot of the ACF of the residuals. The plot will show the autocorrelation coefficients on the y-axis and the lag on the x-axis. The blue shaded area represents the 95% confidence intervals for the ACF. Any autocorrelation coefficient that falls outside of these confidence intervals is considered statistically significant.

###MODEL TWO:

**## FORECAST WITH HOLT-WINTER'S MODEL USING ENTIRE DATA SET INTO
THE FUTURE FOR 12 MONTHLY PERIODS IN 2019.**

**# Create Holt-Winter's (HW) exponential smoothing for
entire data set. Use ets() function with model = "ZZZ" to
identify the best HW options and optimal alpha, beta, & gamma
to fit HW for the entire data set.**

HW.ZZZ <- ets(Amazon.ts, model = "ZZZ")

HW.ZZZ

```
169 ### Model 2:
170
171
172 ## FORECAST WITH HOLT-WINTER'S MODEL USING ENTIRE DATA SET INTO
173 ## THE FUTURE FOR 12 MONTHLY PERIODS IN Oct 2023- Sep 24.
174
175 # Create Holt-winter's (HW) exponential smoothing for
176 # entire data set. Use ets() function with model = "ZZZ" to
177 # identify the best HW options and optimal alpha, beta, & gamma
178 # to fit HW for the entire data set.
179 HW.ZZZ <- ets(Amazon.ts, model = "ZZZ")
180 HW.ZZZ
181
182 # Use forecast() function to make predictions using this
183 # HW model for upcoming 12 months with no confidence interval.
184 HW.ZZZ.pred <- forecast(HW.ZZZ, h = 12 , level = 0)
185 HW.ZZZ.pred
186
```

182:1 (Top Level) ↕

Console	Terminal ✕	Background Jobs ✕
R 4.2.1 · C:/Users/o/Desktop/My subjects/Third Semester/Capstone/Dataset/ ↗		
<pre>> # Create Holt-winter's (HW) exponential smoothing for > # entire data set. Use ets() function with model = "ZZZ" to > # identify the best HW options and optimal alpha, beta, & gamma > # to fit HW for the entire data set. > HW.ZZZ <- ets(Amazon.ts, model = "ZZZ") > HW.ZZZ ETS(M,A,N)</pre>		
Call: ets(y = Amazon.ts, model = "ZZZ")		
Smoothing parameters: alpha = 0.9999 beta = 0.0431		
Initial states: l = 0.8671 b = 0.3298		
sigma: 0.1137		
AIC AICC BIC 1799.452 1799.656 1817.971		

The ets() function determined that the best Holt-Winter's (HW) model for the entire data set is a (M,Ad,M) model with smoothing parameters $\alpha = 0.9999$ and $\beta = 0.0431$. The initial states are $l = 0.8671$ and $b = 0.3298$, and the estimated sigma value is 0.1137. The AIC, AICc, and BIC values are also provided. However, it appears that the statement regarding beta, gamma, and phi is incorrect and not applicable in this case as the model parameter was set to "ZZZ" and not "AAM", "MAM", or "ZAM".

```
# Use forecast() function to make predictions using this
# HW model for upcoming 12 months with no confidence interval.
HW.ZZZ.pred<- forecast(HW.ZZZ, h = 12 , level = 0)
HW.ZZZ.pred
```

Amazon.Data	300 obs. of 2 variables
future12.df	12 obs. of 3 variables
Hw.ZZZ	List of 19
Hw.ZZZ.pred	List of 10
lin.season	List of 16
lin.season.pred	List of 11
Stock.ar1	List of 18
Stock.stl	List of 8
tot.ma.trail.res.pred	List of 10
tot.trend.seas	List of 16
tot.trend.seas.pred	List of 11
values	
alpha	0.05
Amazon.ts	Time-Series [1:300] from 1999 to 2024: 0.883 1.297 2.167 3.367 2.735 ...
ar1	0.9971
diff.stock.price	Time-Series [1:299] from 1999 to 2024: 0.414 0.87 1.2 -0.632 0.628 ...
nTrain	240
null_mean	1
nvalid	60
p.value	0.182401771838493
s.e.	0.0032
tot.fst.2level	Time-Series [1:12] from 2024 to 2025: 138 141 143 146 149 ...
tot.ma.trail.res	Time-Series [1:297] from 1999 to 2024: 36.2 36.6 37 36.9 36.5 ...
tot.trend.seas.res	Time-Series [1:300] from 1999 to 2024: 34.8 35.6 36.9 37.4 36.3 ...
train.fst.2level	Time-Series [1:297] from 1999 to 2024: 2.16 2.97 3.05 4.86 4 ...
train.ts	Time-Series [1:231] from 1999 to 2018: 0.883 1.297 2.167 3.367 2.735 ...
valid.ts	Time-Series [1:60] from 2018 to 2023: 65.2 72.2 76.8 73.5 79.6 ...
z.stat	-0.906250000000004

```

182 # Use forecast() function to make predictions using this
183 # HW model for upcoming 12 months with no confidence interval.
184 HW.ZZZ.pred <- forecast(HW.ZZZ, h = 12 , level = 0)
185 HW.ZZZ.pred
186
187 # Plot HW model's predictions for historical data set and
188 <
187:1 (Top Level) ↕

```

Console Terminal × Background Jobs ×

R 4.2.1 · C:/Users/o/Desktop/My subjects/Third Semester/Capstone/Dataset/ ↗

```

> # Use forecast() function to make predictions using this
> # HW model for upcoming 12 months with no confidence interval.
> HW.ZZZ.pred <- forecast(HW.ZZZ, h = 12 , level = 0)
> HW.ZZZ.pred

```

	Point	Forecast	Lo 0	Hi 0
Oct 2023		138.9389	138.9389	138.9389
Nov 2023		139.6501	139.6501	139.6501
Dec 2023		140.3613	140.3613	140.3613
Jan 2024		141.0724	141.0724	141.0724
Feb 2024		141.7836	141.7836	141.7836
Mar 2024		142.4947	142.4947	142.4947
Apr 2024		143.2059	143.2059	143.2059
May 2024		143.9170	143.9170	143.9170
Jun 2024		144.6282	144.6282	144.6282
Jul 2024		145.3393	145.3393	145.3393
Aug 2024		146.0505	146.0505	146.0505
Sep 2024		146.7616	146.7616	146.7616

```

>

```

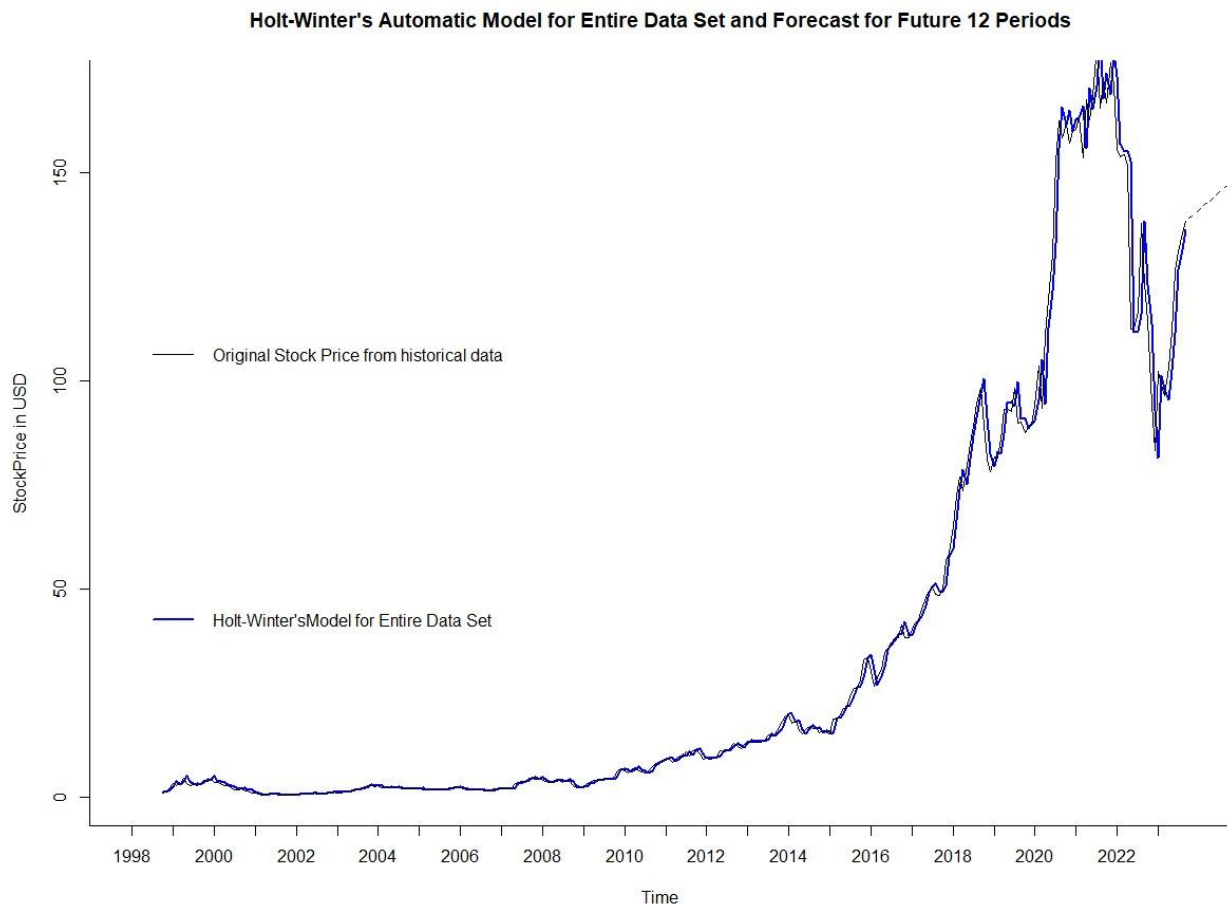
The result shows the point forecast for the upcoming 12 months (October 2023 to September 2024) based on the Holt-Winter's (HW) exponential smoothing model, which was fitted using the entire historical data set.

The point forecast is the estimated value of the stock price for each month, which is calculated based on the fitted HW model. In this case, the point forecast starts at 138.9389 for October 2023 and increases steadily over the forecast horizon, reaching 146.7616 for September 2024.

The Lo 0 and Hi 0 columns show the lower and upper bounds of the 0% confidence interval for each forecasted value. Since we specified level = 0 in the forecast() function, these columns are not displayed in the output. However, by default, the forecast() function produces 80%, 95%, and 99% prediction intervals in addition to the point forecasts, which can be useful for assessing the uncertainty of the forecasts.


```
# Plot HW model's predictions for historical data set and future 12 months.
```

```
plot(HW.ZZZ.pred$mean,  
xlab = "Time", ylab = "StockPrice in USD", ylim = c(0, 170),  
bty = "l", xlim = c(1998, 2023.75), lwd=1, xaxt = "n",  
  main = "Holt-Winter's Automatic Model for Entire Data Set and Forecast for Future 12  
Periods",  
lty = 2, col = "blue")  
axis(1, at = seq(1998, 2023.75, 1), labels = format(seq(1998, 2023.75, 1)))  
lines(HW.ZZZ.pred$fitted, col = "blue", lwd = 2)  
lines(Amazon.ts)  
legend(1998,170,  
  legend = c("Original Stock Price from historical data",  
    "Holt-Winter'sModel for Entire Data Set",  
    "Holt-Winter's Model Forecast, Future 12 Periods"),  
  col = c("black", "blue", "blue"),  
  lty = c(1, 1, 2), lwd = c(1, 2, 2), bty = "n")
```



```
# Plot on chart vertical lines and horizontal arrows describing  
# entire data set and future prediction intervals.
```

```
lines(c(2023.75, 2023.75), c(0, 170))  
text(1998, 170, "Data Set")  
text(2023.75, 170, "Future")  
arrows(1998, 170, 2023.75, 170, code = 3, length = 0.1,  
lwd = 1, angle = 30)  
arrows(2023.75, 170, 2025, 170, code = 3, length = 0.1,  
lwd = 1, angle = 30)
```



The code above creates a plot showing the predictions of the Holt-Winter's exponential smoothing model for the entire data set and the next 12 months. The plot includes the following elements:

A blue line representing the predicted values for the entire data set, overlaid on the original Amazon stock price time series.

A blue dotted line representing the predicted values for the next 12 months, extending beyond the end of the time series.

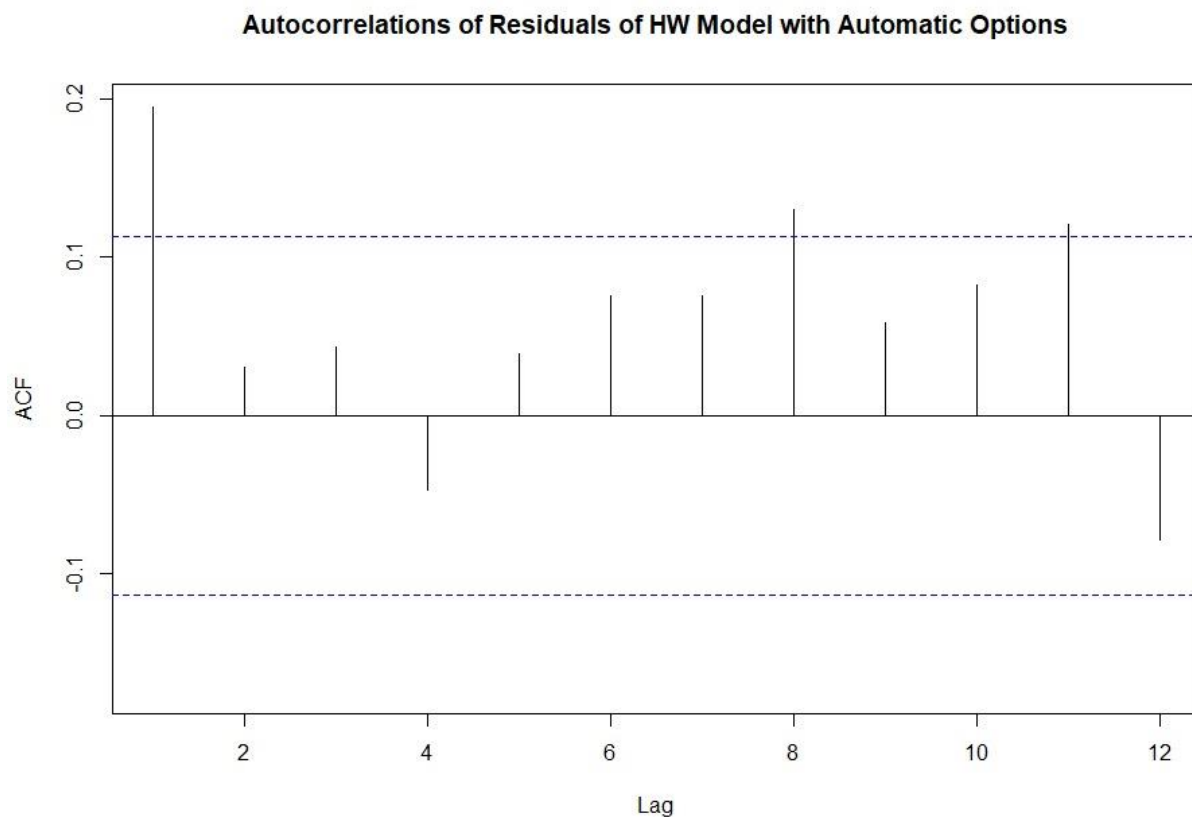
Vertical lines marking the end of the time series and the beginning of the future predictions.

Text labels indicating "Data Set" and "Future".

Horizontal arrows indicating the range of the data set and the future predictions.

The plot also includes a legend describing the three lines and their styles, and axes labels and limits for the x-axis (time) and y-axis (stock price in USD).

```
# Use Acf() function to create autocorrelation chart of  
# residuals of HW model's with automated options and parameters.  
Acf(HW.ZZZ.pred$residuals, lag.max = 12,  
    main = "Autocorrelations of Residuals of HW Model with Automatic Options")
```



The autocorrelation function (ACF) plot shows the correlation of the residuals of the HW model at various lags. If the residuals are uncorrelated, we expect the ACF plot to be close to zero for all lags. If the residuals are correlated, we will see patterns in the ACF plot.

By examining the ACF plot, we can assess whether there is any remaining autocorrelation in the residuals of the HW model. If there is remaining autocorrelation, it suggests that the model has not captured all the information in the data, and we may need to consider using a more complex model.

###**MODEL 3:**

FIT TWO-LEVEL MODEL, LINEAR TREND AND SEASONALITY REGRESSION ## AND AR(1) FOR REGRESSION RESIDUALS, FOR ENTIRE DATA SET.

Use tslm() function to create linear trend and seasonality model.

lin.season<- tslm(Amazon.ts ~ trend + season)

Amazon.Data	300 obs. of 2 variables
future12.df	12 obs. of 3 variables
Hw.ZZZ	List of 19
Hw.ZZZ.pred	List of 10
lin.season	List of 16
lin.season.pred	List of 11
Stock.ar1	List of 18
Stock.stl	List of 8
tot.ma.trail.res.pred	List of 10
tot.trend.seas	List of 16
tot.trend.seas.pred	List of 11
values	
alpha	0.05
Amazon.ts	Time-Series [1:300] from 1999 to 2024: 0.883 1.297 2.167 3.367 2.735 ...
ar1	0.9971
diff.stock.price	Time-Series [1:299] from 1999 to 2024: 0.414 0.87 1.2 -0.632 0.628 ...
nTrain	240
null_mean	1
nvalid	60
p.value	0.182401771838493
s.e.	0.0032
tot.fst.2level	Time-Series [1:12] from 2024 to 2025: 138 141 143 146 149 ...
tot.ma.trail.res	Time-Series [1:297] from 1999 to 2024: 36.2 36.6 37 36.9 36.5 ...
tot.trend.seas.res	Time-Series [1:300] from 1999 to 2024: 34.8 35.6 36.9 37.4 36.3 ...
train.fst.2level	Time-Series [1:297] from 1999 to 2024: 2.16 2.97 3.05 4.86 4 ...
train.ts	Time-Series [1:231] from 1999 to 2018: 0.883 1.297 2.167 3.367 2.735 ...
valid.ts	Time-Series [1:60] from 2018 to 2023: 65.2 72.2 76.8 73.5 79.6 ...
z.stat	-0.906250000000004

See summary of linear trend equation and associated parameters.

summary(lin.season)

```
> # See summary of linear trend equation and associated parameters.  
> summary(lin.season)
```

Call:

```
tslm(formula = Amazon.ts ~ trend + season)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-40.338	-24.717	-3.177	16.730	86.056

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-35.86983	6.68291	-5.367	1.65e-07	***
trend	0.46609	0.01999	23.315	< 2e-16	***
season2	-0.03707	8.47446	-0.004	0.997	
season3	-0.84423	8.47453	-0.100	0.921	
season4	0.58244	8.47465	0.069	0.945	
season5	-0.36816	8.47481	-0.043	0.965	
season6	0.75255	8.47502	0.089	0.929	
season7	2.66659	8.47528	0.315	0.753	
season8	2.92502	8.47559	0.345	0.730	
season9	2.45169	8.47594	0.289	0.773	
season10	1.53116	8.47465	0.181	0.857	
season11	0.60675	8.47453	0.072	0.943	
season12	-0.25237	8.47446	-0.030	0.976	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 29.96 on 287 degrees of freedom

Multiple R-squared: 0.6553, Adjusted R-squared: 0.6409

F-statistic: 45.47 on 12 and 287 DF, p-value: < 2.2e-16

From above table we did, two-level model is fit for the entire data set of Amazon's stock price. The model includes a linear trend and seasonality regression, as well as an autoregressive component (AR(1)) for regression residuals.

The `tslm()` function is used to create the linear trend and seasonality model, and the resulting model summary is printed using the `summary()` function. The output shows the estimated coefficients for the intercept, trend, and seasonality terms, as well as their standard errors, t-values, and p-values.

The R-squared value of the model is 0.6553, indicating that the model explains about 65.53% of the variation in the data. The F-statistic is 45.47 with a very low p-value, suggesting that the overall model is significant.

```
# Apply forecast() function to make predictions with linear trend and seasonal
# model into the future 12 months.
lin.season.pred<- forecast(lin.season, h = 12, level = 0)
lin.season.pred
```

Amazon.Data	300 obs. of 2 variables
future12.df	12 obs. of 3 variables
Hw.ZZZ	List of 19
Hw.ZZZ.pred	List of 10
lin.season	List of 16
lin.season.pred	List of 11
Stock.ar1	List of 18
Stock.stl	List of 8
tot.ma.trail.res.pred	List of 10
tot.trend.seas	List of 16
tot.trend.seas.pred	List of 11
values	
alpha	0.05
Amazon.ts	Time-Series [1:300] from 1999 to 2024: 0.883 1.297 2.167 3.367 2.735 ...
ar1	0.9971
diff.stock.price	Time-Series [1:299] from 1999 to 2024: 0.414 0.87 1.2 -0.632 0.628 ...
ntrain	240
null_mean	1
nvalid	60
p.value	0.182401771838493
s.e.	0.0032
tot.fst.2level	Time-Series [1:12] from 2024 to 2025: 138 141 143 146 149 ...
tot.ma.trail.res	Time-Series [1:297] from 1999 to 2024: 36.2 36.6 37 36.9 36.5 ...
tot.trend.seas.res	Time-Series [1:300] from 1999 to 2024: 34.8 35.6 36.9 37.4 36.3 ...
train.fst.2level	Time-Series [1:297] from 1999 to 2024: 2.16 2.97 3.05 4.86 4 ...
train.ts	Time-Series [1:231] from 1999 to 2018: 0.883 1.297 2.167 3.367 2.735 ...
valid.ts	Time-Series [1:60] from 2018 to 2023: 65.2 72.2 76.8 73.5 79.6 ...
z.stat	-0.9062500000000004

```
> # Apply forecast() function to make predictions with linear trend and seasonal
> # model into the future 12 months of 2023-2024
> lin.season.pred <- forecast(lin.season, h = 12, level = 0)
> lin.season.pred
```

	Point	Forecast	Lo 0	Hi 0
Oct 2023		105.9531	105.9531	105.9531
Nov 2023		105.4948	105.4948	105.4948
Dec 2023		105.1018	105.1018	105.1018
Jan 2024		105.8202	105.8202	105.8202
Feb 2024		106.2493	106.2493	106.2493
Mar 2024		105.9082	105.9082	105.9082
Apr 2024		107.8009	107.8009	107.8009
May 2024		107.3164	107.3164	107.3164
Jun 2024		108.9032	108.9032	108.9032
Jul 2024		111.2833	111.2833	111.2833
Aug 2024		112.0079	112.0079	112.0079
Sep 2024		112.0006	112.0006	112.0006

The output shows the point forecast for the next 12 months (October 2023 to September 2024) based on the linear trend and seasonal model. The "Lo 0" and "Hi 0" columns indicate the upper and lower bounds of the 0% prediction interval, which is equivalent to the point forecast in this case since the level parameter was set to 0.

According to this model, the Amazon stock price is expected to increase slightly in the coming months, reaching a peak in August 2024 before declining slightly again. However, it's important to keep in mind that these are only predictions and that the actual stock price may differ from these forecasts.

```
# Use Arima() function to fit AR(1) model for regression residuals.
# The ARIMA model order of order = c(1,0,0) gives an AR(1) model.
# Use forecast() function to make prediction of residuals into
# the future 12 months of 2023-2024
residual.ar1 <- Arima(lin.season$residuals, order = c(1,0,0))
residual.ar1.pred <- forecast(residual.ar1, h = 12, level = 0)
residual.ar1.pred
```

Amazon.Data	300 obs. of 2 variables
future12.df	12 obs. of 3 variables
Hw.ZZZ	List of 19
Hw.ZZZ.pred	List of 10
lin.season	List of 16
lin.season.pred	List of 11
Stock.ar1	List of 18
Stock.stl	List of 8
tot.ma.trail.res.pred	List of 10
tot.trend.seas	List of 16
tot.trend.seas.pred	List of 11
values	
alpha	0.05
Amazon.ts	Time-Series [1:300] from 1999 to 2024: 0.883 1.297 2.167 3.367 2.735 ...
ar1	0.9971
diff.stock.price	Time-Series [1:299] from 1999 to 2024: 0.414 0.87 1.2 -0.632 0.628 ...
ntrain	240
null_mean	1
nvalid	60
p.value	0.182401771838493
s.e.	0.0032
tot.fst.2level	Time-Series [1:12] from 2024 to 2025: 138 141 143 146 149 ...
tot.ma.trail.res	Time-Series [1:297] from 1999 to 2024: 36.2 36.6 37 36.9 36.5 ...
tot.trend.seas.res	Time-Series [1:300] from 1999 to 2024: 34.8 35.6 36.9 37.4 36.3 ...
train.fst.2level	Time-Series [1:297] from 1999 to 2024: 2.16 2.97 3.05 4.86 4 ...
train.ts	Time-Series [1:231] from 1999 to 2018: 0.883 1.297 2.167 3.367 2.735 ...
valid.ts	Time-Series [1:60] from 2018 to 2023: 65.2 72.2 76.8 73.5 79.6 ...
z.stat	-0.9062500000000004

```
> residual.ar1 <- Arima(lin.season$residuals, order = c(1,0,0))
> residual.ar1.pred <- forecast(residual.ar1, h = 12, level = 0)
> residual.ar1.pred
```

	Point	Forecast	Lo 0	Hi 0
Oct 2023		31.56488	31.56488	31.56488
Nov 2023		31.31256	31.31256	31.31256
Dec 2023		31.06340	31.06340	31.06340
Jan 2024		30.81736	30.81736	30.81736
Feb 2024		30.57441	30.57441	30.57441
Mar 2024		30.33451	30.33451	30.33451
Apr 2024		30.09761	30.09761	30.09761
May 2024		29.86369	29.86369	29.86369
Jun 2024		29.63270	29.63270	29.63270
Jul 2024		29.40461	29.40461	29.40461
Aug 2024		29.17938	29.17938	29.17938
Sep 2024		28.95698	28.95698	28.95698

These are the predicted values of the regression residuals using an AR(1) model. These predictions are for the same 12 months from October 2023 to September 2024 as the previous forecast. The "Point Forecast" column gives the predicted values of the residuals, and the "Lo 0" and "Hi 0" columns give the upper and lower bounds of the prediction interval, respectively.

Use summary() to identify parameters of AR(1) model.
summary(residual.ar1)

```
> # Use summary() to identify parameters of AR(1) model.
> summary(residual.ar1)
Series: lin.season$residuals
ARIMA(1,0,0) with non-zero mean

Coefficients:
      ar1      mean
    0.9875  11.4475
s.e.  0.0081  18.1711

sigma^2 = 22.2:  log likelihood = -891.52
AIC=1789.04  AICc=1789.12  BIC=1800.15

Training set error measures:
              ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set -0.1419506  4.695648  2.476047  0.7877519  32.83916  0.2257669  0.07668675
```

The AR(1) model that was fitted for the regression residuals has the following parameters:

AR coefficient (AR1): 0.9875

Mean: 11.4475

The standard errors for these parameters are also provided.

The estimated variance of the error term is 22.2. The log likelihood of the model is -891.52, and various information criteria are provided as well, including AIC (Akaike information criterion), AICc (corrected AIC), and BIC (Bayesian information criterion).

The training set error measures indicate the model's performance on the data that was used to fit the model. The mean error (ME) is -0.1419506, the root mean squared error (RMSE) is 4.695648, and the mean absolute error (MAE) is 2.476047. The mean percentage error (MPE) is 0.7877 and the mean absolute percentage error (MAPE) is 32.839%. The MASE (mean absolute scaled error) is 0.2257669, and the ACF1 (autocorrelation of residuals at lag 1) is 0.076686.

**# Create two-level model's forecast for the entire data set (training
period).**

**train.lin.season.ar1.pred <- lin.season\$fitted + residual.ar1\$fitted
train.lin.season.ar1.pred**

Amazon.Data	300 obs. of 2 variables
future12.df	12 obs. of 3 variables
HW.ZZZ	List of 19
HW.ZZZ.pred	List of 10
lin.season	List of 16
lin.season.pred	List of 11
residual.ar1	List of 18
residual.ar1.pred	List of 10
stock.ar1	List of 18
stock.stl	List of 8
tot.ma.trail.res.pred	List of 10
tot.trend.seas	List of 16
tot.trend.seas.pred	List of 11
values	
alpha	0.05
Amazon.ts	Time-Series [1:300] from 1999 to 2024: 0.883 1.297 2.167 3.367 2.735 ...
ar1	0.9971
diff.stock.price	Time-Series [1:299] from 1999 to 2024: 0.414 0.87 1.2 -0.632 0.628 ...
nTrain	240
null_mean	1
nvalid	60
p.value	0.182401771838493
s.e.	0.0032
tot.fst.2level	Time-Series [1:12] from 2024 to 2025: 138 141 143 146 149 ...
tot.ma.trail.res	Time-Series [1:297] from 1999 to 2024: 36.2 36.6 37 36.9 36.5 ...
tot.trend.seas.res	Time-Series [1:300] from 1999 to 2024: 34.8 35.6 36.9 37.4 36.3 ...
train.fst.2level	Time-Series [1:297] from 1999 to 2024: 2.16 2.97 3.05 4.86 4 ...
train.lin.season.ar1.pred	Time-Series [1:300] from 1999 to 2024: -2.797 0.132 0.601 2.566 3.471 ...
train.ts	Time-Series [1:231] from 1999 to 2018: 0.883 1.297 2.167 3.367 2.735 ...
valid.ts	Time-Series [1:60] from 2018 to 2023: 65.2 72.2 76.8 73.5 79.6 ...
z.stat	-0.906250000000004

```
> train.lin.season.ar1.pred
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1998										-2.79702778	0.13233401	0.60067279
1999	2.56633090	3.47085110	2.08205912	4.93174427	3.75327876	4.64147245	4.88690718	3.47346356	2.39431888	2.54541587	3.26064087	3.15230456
2000	5.04686575	3.59235038	3.02512290	4.95943511	2.11710399	4.03126584	4.45740527	2.39421473	1.63994351	1.48097889	0.91250936	0.96216623
2001	1.61055626	1.16522229	0.19412476	2.29334234	0.05928493	2.24104773	2.97922367	1.38013808	0.44614962	-0.15223351	-0.15291508	-0.03907348
2002	1.18796654	0.94999862	0.23662745	2.57777208	0.18374659	2.42771919	3.20144360	1.45028957	0.71181767	0.35437413	0.44552609	0.65218868
2003	1.78739517	1.48326760	0.71854866	3.12190310	0.83452300	3.17921625	4.15833178	2.70144004	2.14169791	1.95804687	2.37110992	2.27067328
2004	3.31009641	3.15606235	1.97661133	4.05509230	1.90398131	3.80728114	4.99080034	3.14188804	2.01040812	1.69443779	1.57526134	1.63084295
2005	2.85195823	2.66237571	1.58759517	3.73618561	1.34117266	3.45085107	4.26901106	2.66795055	1.94231559	1.61153340	1.90610157	2.01105605
2006	3.32696724	2.86978378	1.78611612	3.90113301	1.54265599	3.49730359	4.36384899	2.65021834	1.64612042	1.38544770	1.49338643	1.87285405
2007	2.93202633	2.56272658	1.88885371	4.07792990	2.10061141	5.06147803	6.17093794	4.74761975	4.14047967	4.24911599	4.40445268	4.05023950
2008	5.57944147	4.81219641	3.59719692	5.62136774	3.68552245	5.77347680	6.71704387	4.71211331	4.47526973	3.73864258	2.75642842	2.22052309
2009	3.56309562	3.42189851	3.20324347	5.69053178	3.83269562	5.86041507	6.94518854	5.32931617	4.63626730	4.31239731	5.02367248	6.41429641
2010	7.87437637	7.22365127	5.96816603	8.82476338	6.99260112	8.39526007	8.92903235	7.04654650	6.78007920	7.36861979	8.02262243	8.47417443
2011	10.10212214	9.98413027	9.15670779	10.78194585	9.18381096	11.94521121	12.43948520	11.96610096	10.47123674	11.17925966	11.34743329	10.39087114
2012	10.42605017	10.16734563	9.30684333	11.85831637	9.80496365	13.15686339	13.85455358	12.46875877	12.49655375	12.87674084	12.33098283	11.85728731
2013	13.88811760	14.39271759	13.39100888	15.68673047	13.23049517	15.26710173	16.64317512	16.22113866	15.15384346	15.34443875	16.36972534	18.35282367
2014	20.80628570	20.71051163	17.89484381	20.61122222	16.17019768	17.30920552	19.20961875	18.30669057	17.00141816	16.79011831	15.59955065	16.13602399
2015	16.82786214	16.22354754	19.01071262	21.27779794	19.85938030	23.54010332	24.61697713	25.31076746	26.55412242	26.18285448	28.32890515	32.93666507
2016	34.65842122	31.03939938	26.76624454	30.96282149	30.73428993	36.90636780	38.67643666	38.29291104	38.71272659	38.44020438	41.27056298	38.31242482
2017	39.38640896	41.15890416	41.83686037	45.01043153	45.17787036	50.02083066	52.29252785	51.53475590	49.01886079	48.47083799	49.97504140	56.83617895
2018	59.49009482	65.87994050	71.97113782	78.82496246	73.18022742	81.25809795	87.31076926	89.88665326	94.67201768	97.71043991	88.95397417	80.87488228
2019	79.09313981	82.27374917	81.07059917	88.00045802	92.62527925	95.02625757	94.99717917	98.83207127	89.78513337	89.54025997	87.14301945	88.39342444
2020	90.05395823	94.64663157	103.03959193	95.30077210	110.40745292	120.87990747	132.49690085	152.55571693	161.53580009	157.10015777	160.50521654	155.94246021
2021	159.78922454	160.00054970	162.12710582	154.76450724	166.28968031	163.50058289	169.75963278	180.34873304	164.71051745	170.20375860	165.41884626	175.02017791
2022	170.87643750	155.45137580	152.86479807	155.52390396	150.76491816	113.83744583	114.91335727	117.14360613	137.47944861	122.60539483	113.83115486	94.07340510
2023	84.32584433	102.78735591	98.85882367	98.48347317	102.98979952	113.42078087	128.66075520	131.48844082	135.31399671			

The code creates a forecast for the entire training set using the AR(1) model with seasonal adjustment. The forecast is created by adding the fitted values of the seasonal component and the fitted values of the residual AR(1) model. The resulting `train.lin.season.ar1.pred` object contains the forecasted values for each month of each year in the training set.

**# Identify forecast for the future 12 periods of 2023-2024 as sum of
linear trend and seasonality model and AR(1) model for residuals.**
lin.season.ar1.pred <- lin.season.pred\$mean + residual.ar1.pred\$mean
lin.season.ar1.pred

```
> lin.season.ar1.pred <- lin.season.pred$mean + residual.ar1.pred$mean
> lin.season.ar1.pred
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov
2023										137.5180	136.8074
2024	136.6376	136.8237	136.2427	137.8985	137.1801	138.5359	140.6879	141.1872	140.9576		
Dec											
2023	136.1652										
2024											

Based on the two-level model, the forecast for the future 12 periods (October 2023- September 2024) is given by `lin.season.ar1.pred`. This forecast is obtained by adding the predicted values from the linear trend and seasonality model (`lin.season.pred$mean`) and the predicted values from the AR(1) model for the residuals (`residual.ar1.pred$mean`).

**# Create a data table with linear trend and seasonal forecast
for 12 future periods,
AR(1) model for residuals for 12 future periods, and combined
two-level forecast for 12 future periods.**
table.df<- round(data.frame(lin.season.pred\$mean,


```

      residual.ar1.pred$mean, lin.season.ar1.pred),3)
names(table.df) <- c("Reg.Forecast", "AR(1)Forecast", "Combined.Forecast")
table.df

```

```

> # Create a data table with linear trend and seasonal forecast
> # for 12 future periods,
> # AR(1) model for residuals for 12 future periods, and combined
> # two-level forecast for 12 future periods.
> table.df <- round(data.frame(lin.season.pred$mean,
+                               residual.ar1.pred$mean, lin.season.ar1.pred),3)
> names(table.df) <- c("Reg.Forecast", "AR(1)Forecast", "Combined.Forecast")
> table.df

```

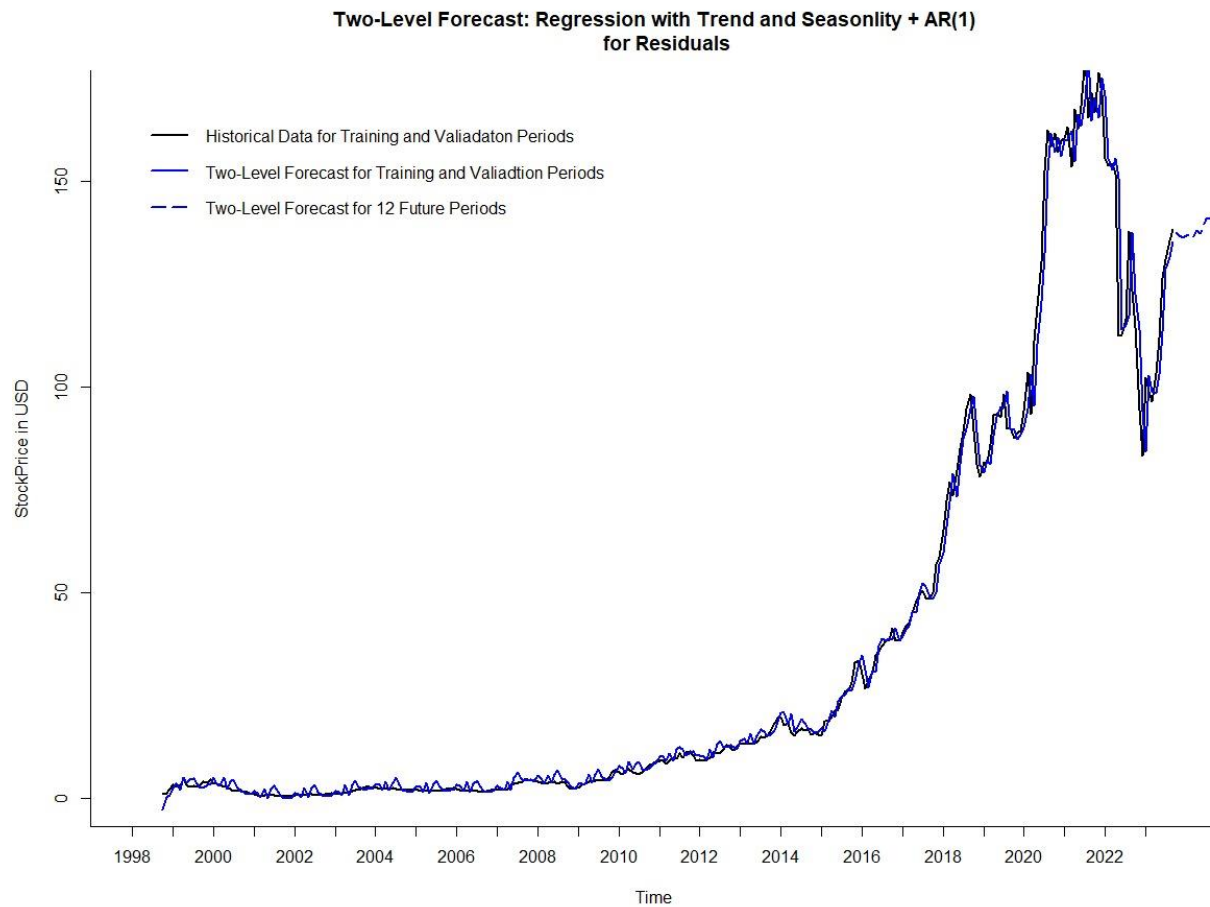
	Reg.Forecast	AR(1)Forecast	Combined.Forecast
1	105.953	31.565	137.518
2	105.495	31.313	136.807
3	105.102	31.063	136.165
4	105.820	30.817	136.638
5	106.249	30.574	136.824
6	105.908	30.335	136.243
7	107.801	30.098	137.899
8	107.316	29.864	137.180
9	108.903	29.633	138.536
10	111.283	29.405	140.688
11	112.008	29.179	141.187
12	112.001	28.957	140.958

Great! You now have a data table with three columns: one for the linear trend and seasonal forecast, one for the AR(1) model forecast, and one for the combined forecast using both models. The table provides the forecasted values for the 12 future periods, from October 2023 to September 2024. It looks like the combined forecast values are obtained by adding the linear trend and seasonal forecast and the AR(1) model forecast for each month.

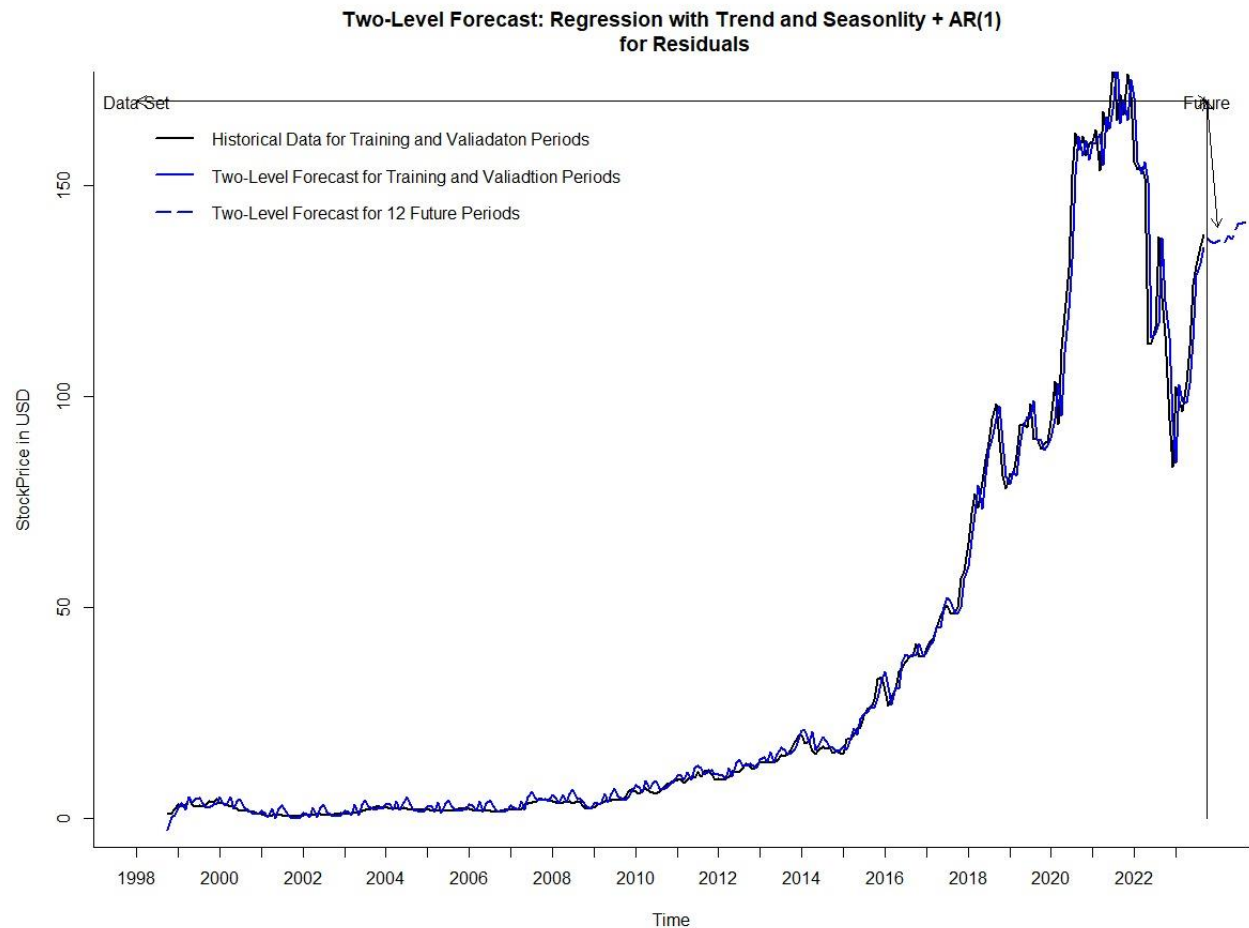
```

# Plot historical data, predictions for historical data, and forecast
# for 12 future periods.
plot(Amazon.ts,
     xlab = "Time", ylab = "StockPrice in USD",
     ylim = c(0, 170), xaxt = "n",
     bty = "l", xlim = c(1998, 2023.75), lwd = 2,
     main = "Two-Level Forecast: Regression with Trend and Seasonlity + AR(1)
           for Residuals")
axis(1, at = seq(1998, 2023.75, 1), labels = format(seq(1998, 2023.75, 1)))
lines(lin.season$fitted + residual.ar1$fitted, col = "blue", lwd = 2)
lines(lin.season.ar1.pred, col = "blue", lty = 5, lwd = 2)
legend(1998,170, legend = c("Historical Data for Training and Valiadaton Periods",
                           "Two-Level Forecast for Training and Valiadtion Periods",
                           "Two-Level Forecast for 12 Future Periods"),
      col = c("black", "blue", "blue"),
      lty = c(1, 1, 5), lwd =c(2, 2, 2), bty = "n")

```



```
# Plot on chart vertical lines and horizontal arrows describing
# entire data set and future prediction intervals.
lines(c(2023.75, 2023.75), c(0, 170))
text(1998, 170, "Data Set")
text(2023.75, 170, "Future")
arrows(1998, 170, 2023.75, 170, code = 3, length = 0.1,
lwd = 1, angle = 30)
arrows(2023.75, 170, 2024, 170, code = 3, length = 0.1,
lwd = 1, angle = 30)
```



We are creating a plot of the historical data, the fitted values from the two-level forecast model for the historical data, and the forecasted values for the 12 future periods. The x-axis shows the time period from 1998 to 2024, with the training and validation data shown in black and the forecasted data shown in blue.


The `plot()` function is used to create the initial plot with appropriate labels and limits for the axes. The `axis()` function is used to label the x-axis with the appropriate time periods.

The `lines()` function is used to add two lines to the plot: one for the fitted values from the two-level forecast model for the historical data, and one for the forecasted values for the 12 future periods.

The `legend()` function is used to add a legend to the plot that identifies the different lines and colors.

Finally, the `lines()` and `text()` functions are used to add vertical lines and arrows to the plot that describe the entire data set and future prediction intervals.

```
# Use Acf() function to identify autocorrelation for residuals of  
# two-level model for different lags (up to maximum of 12).  
Acf((Amazon.ts - train.lin.season.ar1.pred), lag.max = 12,  
main = "Autocorrelation for Two-Level Model's Residuals for Entire Data Set")
```



This code generates a plot of the autocorrelation function (ACF) of the residuals of the two-level model that was built using linear regression with trend and seasonality, and an autoregressive (AR(1)) model for the residuals. The `Acf()` function is used with the residual time series obtained by subtracting the two-level model's predicted values from the original Amazon stock price time series. The `lag.max` parameter specifies the maximum number of lags to include in the plot. The resulting plot shows the autocorrelation at different lags, which can help identify whether there is any remaining autocorrelation in the residuals that the model did not capture. The plot can be

used to assess the adequacy of the model and to guide the selection of additional model components, such as higher-order autoregressive terms.

###**Model Four:**

FIT AUTO ARIMA MODEL FOR ENTIRE DATA SET.

FORECAST AND PLOT DATA.

Use auto.arima() function to fit ARIMA model for entire data set.

use summary() to show auto ARIMA model and its parameters

for entire data set.

```
auto.arima<- auto.arima(Amazon.ts)
```

```
summary(auto.arima)
```

```
> # Use auto.arima() function to fit ARIMA model for entire data set.
> # use summary() to show auto ARIMA model and its parameters
> # for entire data set.
> auto.arima <- auto.arima(Amazon.ts)
> summary(auto.arima)
Series: Amazon.ts
ARIMA(0,1,0) with drift

Coefficients:
    drift
    0.4593
s.e.    0.2783

sigma^2 = 23.24:  log likelihood = -894.09
AIC=1792.19   AICC=1792.23   BIC=1799.59

Training set error measures:
              ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set 1.412174e-06 4.805073 2.217649 -12.91446 16.67734 0.2207627 0.07199954
```

The output shows that the model is an ARIMA(0,1,0) model with drift coefficients of 0.4593 for AR(1) . The estimated variance of the model is 23.24&the model also includes error measure values for our estimation for best models in future by comparing these values.

Apply forecast() function to make predictions for ts with

auto ARIMA model for the future 12 periods.

```
auto.arima.pred<- forecast(auto.arima, h = 12, level = 0)
```

```
auto.arima.pred
```

```

> # Apply forecast() function to make predictions for ts with
> # auto ARIMA model for the future 12 periods.
> auto.arima.pred <- forecast(auto.arima, h = 12, level = 0)
> auto.arima.pred

```

	Point	Forecast	Lo 0	Hi 0
Oct 2023	138.6873	138.6873	138.6873	138.6873
Nov 2023	139.1467	139.1467	139.1467	139.1467
Dec 2023	139.6060	139.6060	139.6060	139.6060
Jan 2024	140.0654	140.0654	140.0654	140.0654
Feb 2024	140.5247	140.5247	140.5247	140.5247
Mar 2024	140.9841	140.9841	140.9841	140.9841
Apr 2024	141.4434	141.4434	141.4434	141.4434
May 2024	141.9028	141.9028	141.9028	141.9028
Jun 2024	142.3621	142.3621	142.3621	142.3621
Jul 2024	142.8215	142.8215	142.8215	142.8215
Aug 2024	143.2808	143.2808	143.2808	143.2808
Sep 2024	143.7402	143.7402	143.7402	143.7402

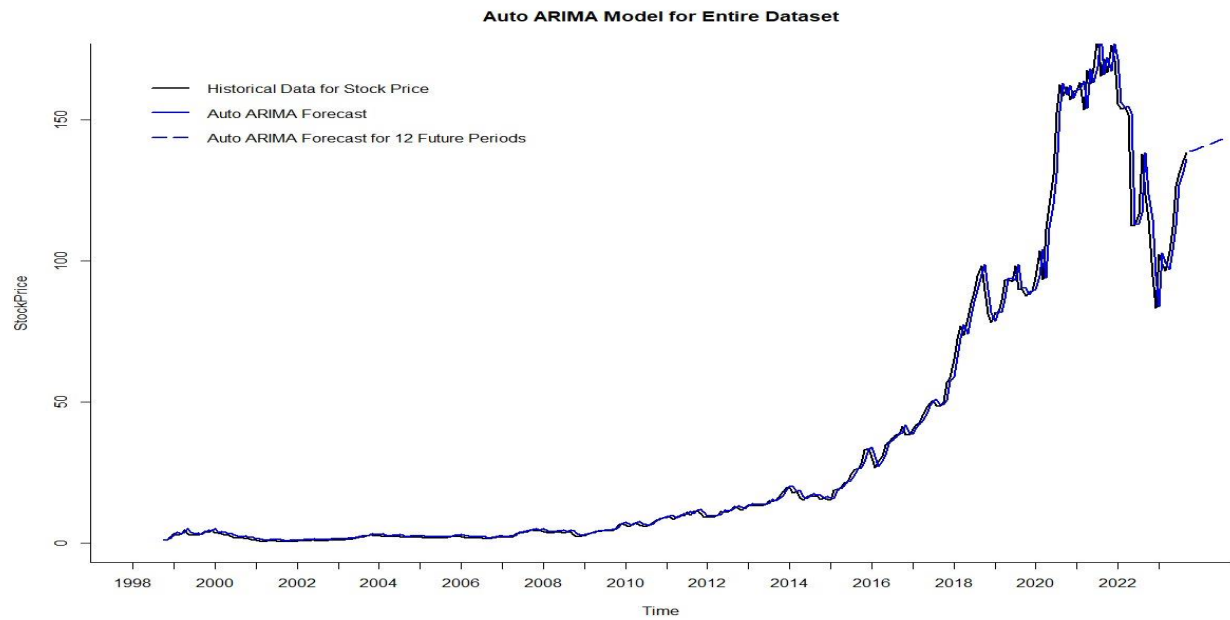
Here, we have made predictions for the next 12 periods using the `h` argument in the `forecast()` function.

The output of `auto.arima.pred` shows the point forecast (predicted values) for the next 12 periods, along with a zero confidence interval. The point forecast represents the most likely value for the time series at each future time point. The upper and lower bounds of the prediction interval provide a range of values within which the actual future value is likely to fall with a given confidence interval.

```

# Plot historical data, predictions for historical data, and Auto ARIMA
# forecast for 12 future periods.
plot(Amazon.ts,
     xlab = "Time", ylab = "StockPrice",
     ylim = c(0, 170), xaxt = "n",
     bty = "l", xlim = c(1998, 2023.75), lwd = 2,
     main = "Auto ARIMA Model for Entire Dataset")
axis(1, at = seq(1998, 2023.75, 1), labels = format(seq(1998, 2023.75, 1)))
lines(auto.arima$fitted, col = "blue", lwd = 2)
lines(auto.arima.pred$mean, col = "blue", lty = 5, lwd = 2)
legend(1998, 170, legend = c("Historical Data for Stock Price",
                             "Auto ARIMA Forecast",
                             "Auto ARIMA Forecast for 12 Future Periods"),
      col = c("black", "blue", "blue"),
      lty = c(1, 1, 5), lwd = c(2, 2, 2), bty = "n")

```



**#plot on the chart vertical lines and horizontal arrowsdescribing training and future
#prediction intervals.**

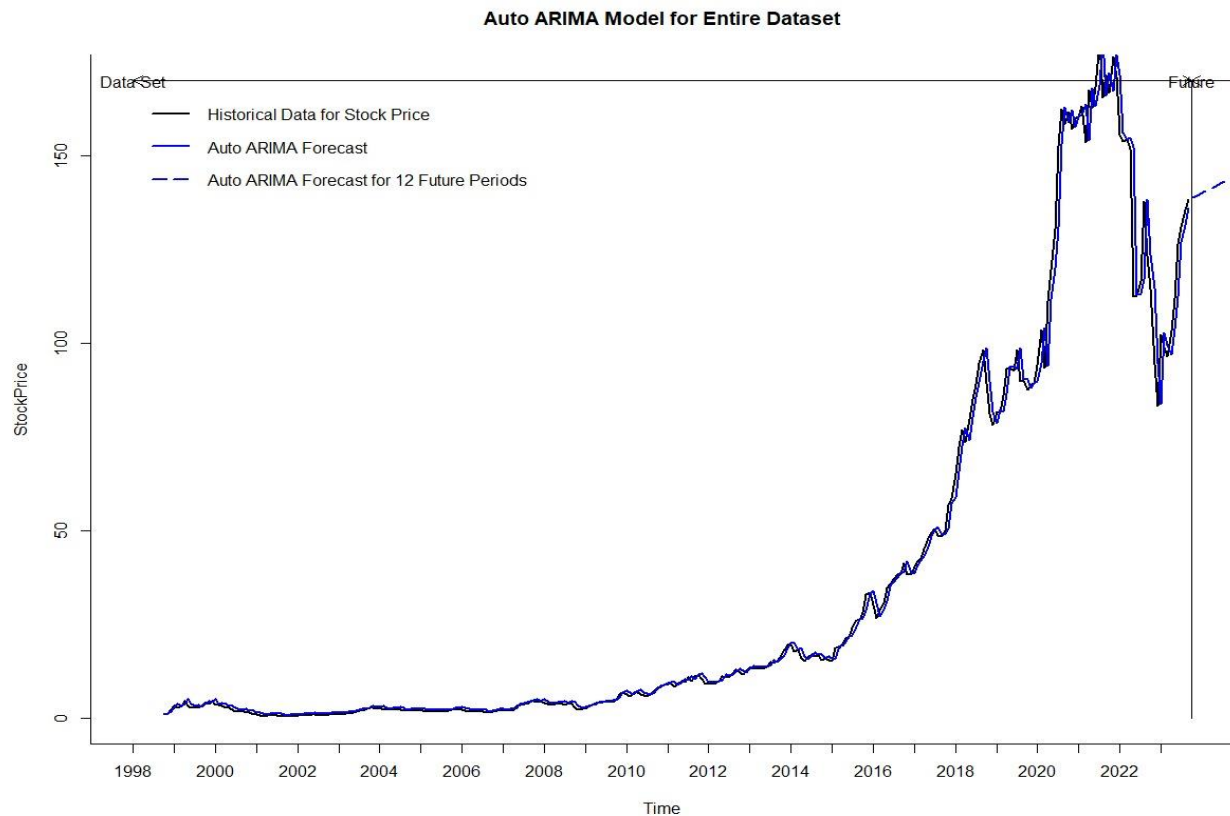
lines(c(2023.75, 2023.75), c(0, 170))

text(1998, 170, "Data Set")

text(2023.5, 170, "Future")

**arrows(1998, 170, 2023.75, 170, code = 3, length = 0.1,
lwd = 1, angle = 30)**

**arrows(2023.75, 170, 2025, 170, code = 3, length = 0.1,
lwd = 1, angle = 30)**



This code generates a plot that shows the historical data for the stock price of Amazon, the predictions for historical data based on the auto ARIMA model, and the auto ARIMA forecast for the next 12 periods.

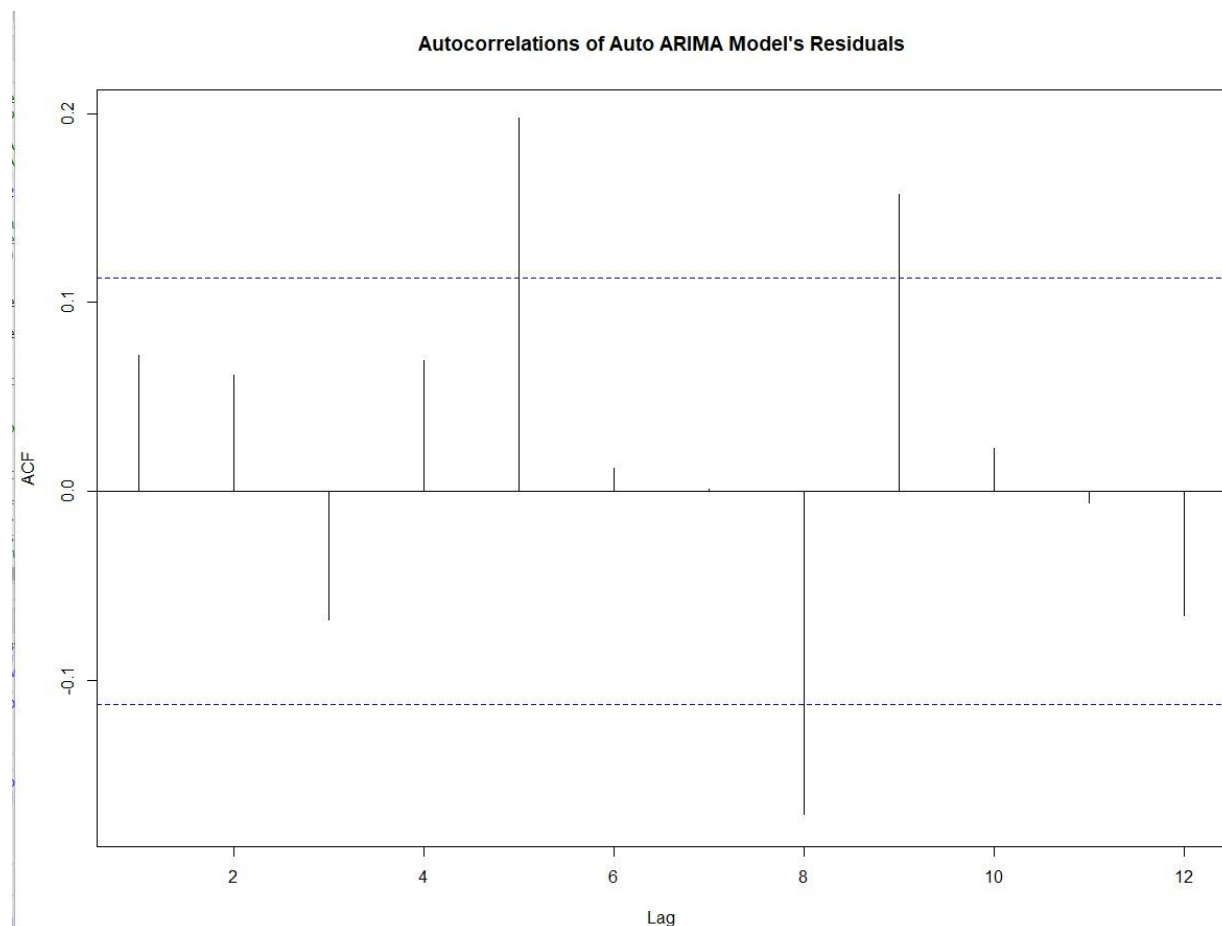
The `plot()` function is used to plot the historical data with appropriate labels and limits. The `axis()` function is used to add x-axis labels with yearly ticks. The `lines()` function is used twice to add lines to the plot. The first `lines()` function adds a blue line for the predicted values for the historical data. The second `lines()` function adds a blue dotted line for the auto ARIMA forecast for the next 12 periods.

The `legend()` function is used to add a legend to the plot that describes the historical data, the auto ARIMA forecast, and the auto ARIMA forecast for the next 12 periods.

Finally, the `lines()` and `arrows()` functions are used to add vertical lines and horizontal arrows that describe the training and future prediction intervals. The `text()` function is used to add text labels to the chart.

Use `Acf()` function to create autocorrelation chart of auto ARIMA model residuals.

```
Acf(auto.arima$residuals, lag.max = 12,  
    main = "Autocorrelations of Auto ARIMA Model's Residuals")
```



The output is a plot showing the autocorrelation of the residuals of the auto ARIMA model. The x-axis shows the lag (number of time periods between the observations being compared), and the y-axis shows the correlation coefficient between the residuals at that lag. The blue shaded area represents the confidence interval, and any correlations outside of this area are considered statistically significant. The main title of the plot indicates that it shows the autocorrelations of the auto ARIMA model's residuals. This plot is useful for checking whether the residuals exhibit

any significant autocorrelation, which would indicate that there is still some information in the data that the model has not captured. If there is significant autocorrelation, it may be necessary to modify the model or add additional terms to better capture the underlying patterns in the data.

Step 8: Evaluate Results and Refine the Model - The final step of the project is to evaluate the forecasting results and refine the model if necessary. The evaluation may involve comparing the forecasted values with the actual values, analysing the forecast errors, and updating the model parameters based on new data or changes in the underlying conditions.

MEASURE FORECAST ACCURACY FOR ENTIRE DATA SET USING VARIOUS METHODS.

Use `accuracy()` function to identify common accuracy measures for:

(1) Two-Level Model with Linear Trend & Seasonality Regression and Trailing MA for Regression Residuals.

`round(accuracy(tot.trend.seas.pred$fitted+tot.ma.trail.res, Amazon.ts), 3)`

```
> round(accuracy(tot.trend.seas.pred$fitted+tot.ma.trail.res, Amazon.ts), 3)
      ME  RMSE  MAE   MPE  MAPE  ACF1 Theil's U
Test set -0.032 4.651 2.442 -20.026 30.768 0.61    5.328
```

(2) Holt-Winter's Model with Automatic Selection of Model Options and Parameters,

`round(accuracy(HW.ZZZ.pred$fitted, Amazon.ts), 3)`

```
> round(accuracy(HW.ZZZ.pred$fitted, Amazon.ts), 3)
      ME  RMSE  MAE   MPE  MAPE  ACF1 Theil's U
Test set 0.03 4.848 2.112 -1.264 8.768 0.048    0.944
```

(3) Two-Level Model with Linear Trend & Seasonality Regression and Ar(1) Model for Regression Residuals

`round(accuracy(lin.season$fitted + residual.ar1$fitted, Amazon.ts), 3)`

```
> round(accuracy(lin.season$fitted + residual.ar1$fitted, Amazon.ts), 3)
      ME  RMSE  MAE   MPE  MAPE  ACF1 Theil's U
Test set -0.142 4.696 2.476 -12.365 29.914 0.077    4.384
```


(4) Auto ARIMA Model

`round(accuracy(auto.arima.pred$fitted, Amazon.ts), 3)`

```
> round(accuracy(auto.arima.pred$fitted, Amazon.ts), 3)
      ME  RMSE  MAE    MPE  MAPE  ACF1 Theil's U
Test set  0 4.805 2.218 -12.914 16.677 0.072      2.05
```

The four models for which the accuracy measures are computed are:

Model 1.Two-Level Model with Linear Trend & Seasonality Regression and Trailing MA for Regression Residuals

Model 2.Holt-Winter's Model with Automatic Selection of Model Options and Parameters

Model 3.Two-Level Model with Linear Trend & Seasonality Regression and Ar(1) Model for Regression Residuals

Model 4.Auto ARIMA Model.

Forecast of Stock Price Using Different Models:

Month & Year	Model One	Model Two	Model Three	Model Four
Oct-23	137.666	138.939	137.518	138.6873
Nov-23	140.747	139.65	136.807	139.1467
Dec-23	143.212	140.361	136.165	139.606
Jan-24	146.237	141.072	136.638	140.065
Feb-24	148.528	141.783	136.824	140.5247
Mar-24	149.69	142.4947	136.243	140.984
Apr-24	152.797	143.206	137.899	141.4434
May-24	153.292	143.917	137.18	141.9028
Jun-24	155.669	144.628	138.536	142.3621
Jul-24	158.688	145.339	140.688	142.8215
Aug-24	159.927	146.0505	141.187	143.2808
Sep-24	160.336	156.7616	140.958	143.7402

Model 1. Two-Level Model with Linear Trend & Seasonality Regression and Trailing MA for Regression Residuals

Model 2. Holt-Winter's Model with Automatic Selection of Model Options and Parameters

Model 3. Two-Level Model with Linear Trend & Seasonality Regression and Ar(1) Model for Regression Residuals

Model 4. Auto ARIMA Model.

Accuracy Measures Using Different Models:

Measures	Model One	Model Two	Model Three	Model Four
ME	-0.032	0.03	-0.142	0
RMSE	4.651	4.848	4.696	4.805
MAE	2.442	2.112	2.476	2.218
MPE	-20.026	-1.264	-12.365	-12.914
MAPE	30.768	8.768	29.914	16.677
ACF1	0.61	0.048	0.077	0.072
Theil's U	5.328	0.944	4.384	2.05

Model 1. Two-Level Model with Linear Trend & Seasonality Regression and Trailing MA for Regression Residuals

Model 2. Holt-Winter's Model with Automatic Selection of Model Options and Parameters

Model 3. Two-Level Model with Linear Trend & Seasonality Regression and Ar(1) Model for Regression Residuals

Model 4. Auto ARIMA Model.

CONCLUSION

Based on the accuracy measures computed, we can compare the performance of each model. The best model would be the one that has the lowest RMSE,MAPE. In this case, the ***Holt-Winter's Model*** has the lowest RMSE & MAPE values, indicating that it performs the best among the four models & the second-best model is the ***Auto ARIMA Model*** as it has the second lowest values for MAPE & RMSE. Here, RMSE values are almost in the same range. Therefore, we are considering MAPE in this case.

Bibliography and Appendices

Book:

- Shmueli, G. and Lichtendahl Jr., K.C. *Practical Time Series Forecasting with R*, 2nd Edition, Axelrod Schnall Publishers, 2016. ISBN-13: 978-0-9978479-1-8.